

Anexo III. Código y resultados para Aprendizaje Automático

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from datetime import datetime
from math import sqrt, exp
from statistics import mean

from statsmodels.api import ProbPlot
from statsmodels.graphics import tsaplots
from statsmodels.stats.stattools import jarque_bera as jb
from statsmodels.stats.diagnostic import acorr_ljungbox as lb
from statsmodels.tsa.seasonal import seasonal_decompose

from scipy.stats import kstest, boxcox, gaussian_kde, norm

from pmdarima.arima import ndiffs
from pmdarima.arima import nsdiffs

from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import mean_absolute_error as mae

from tstoolbox import *

import warnings
warnings.filterwarnings("ignore")

In [2]: sns.set(rc = {"figure.figsize": (10,4), "axes.facecolor": "#eeeef4"})
my_palette = sns.color_palette(sns.diverging_palette(255,
                                                    133,
                                                    l=40,
                                                    n=4,
```

```

sns.set_palette(my_palette)

In [3]: ts_df = pd.read_csv('wolf_river.csv', sep = ';')
        ts_df['index'] = pd.to_datetime(ts_df['month'])
        ts_df.set_index('index', inplace = True)
        ts_df.drop(["month"], axis = 1, inplace = True) # Serie como pd.DataFrame
        ts = ts_df.iloc[:,0] # Serie como pd.Series

In [4]: type(ts_df)


Out[4]: pandas.core.frame.DataFrame

In [5]: type(ts)

Out[5]: pandas.core.series.Series

In [6]: tsplot(ts)
        plt.savefig("river.png", dpi = 400)
        plt.show()

```




División temporal en período de modelización y período de predicción

```

In [7]: train = ts_split(ts_df)[0] # Período de ajuste de modelos
        test  = ts_split(ts)[1]    # Período de generación de predicciones

In [8]: train_test_plot(train, test)
        plt.savefig("river_split.png", dpi = 400)
        plt.show()

```



Fechas destacables

```
In [9]: start_date = train.index[0]
        split_date = test.index[0] # Primera observación a predecir
        end_date   = test.index[-1]
```

Preprocesamientos para la serie

```
In [10]: boxcox_trans = boxcox(train)
         bc_param = boxcox_trans[1][0]
         bc_train = ((train**bc_param)-1)/bc_param

In [11]: ts_components = seasonal_decompose(train, "add")
         seasonal = ts_components.seasonal
         deseas = train - seasonal

         seasonality_forecast = snaive(seasonal, test)

In [12]: diff_train = train.diff(1).iloc[1:]

In [13]: bc_ts_components = seasonal_decompose(bc_train, "add")
         bc_seasonal = bc_ts_components.seasonal
         bc_deseas = bc_train - bc_seasonal

         bc_seasonality_forecast = snaive(bc_seasonal, test)

In [14]: diff_bc_train = bc_train.diff(1).iloc[1:]

In [15]: diff_deseas = deseas.diff(1).iloc[1:]

In [16]: bc_diff_deseas = bc_deseas.diff(1).iloc[1:]
```

Fijación de semilla para reproducibilidad

```
In [17]: seed = 1
```

0.0.1 Serie temporal original

```
In [18]: lag_df = series_to_supervised(train, n_in=12)
lag_df.head(5)
```

```
Out[18]:
```

	var1(t-12)	var1(t-11)	var1(t-10)	var1(t-9)	var1(t-8)	\
index						
1915-01-01	55.275	49.951	42.617	26.901	22.654	
1915-02-01	49.951	42.617	26.901	22.654	35.793	
1915-03-01	42.617	26.901	22.654	35.793	66.063	
1915-04-01	26.901	22.654	35.793	66.063	83.903	
1915-05-01	22.654	35.793	66.063	83.903	113.636	

	var1(t-7)	var1(t-6)	var1(t-5)	var1(t-4)	var1(t-3)	var1(t-2)	\
index							
1915-01-01	35.793	66.063	83.903	113.636	60.655	36.557	
1915-02-01	66.063	83.903	113.636	60.655	36.557	51.537	
1915-03-01	83.903	113.636	60.655	36.557	51.537	41.739	
1915-04-01	113.636	60.655	36.557	51.537	41.739	37.492	
1915-05-01	60.655	36.557	51.537	41.739	37.492	31.998	

	var1(t-1)	var1(t)
index		
1915-01-01	51.537	41.739
1915-02-01	41.739	37.492
1915-03-01	37.492	31.998
1915-04-01	31.998	22.823
1915-05-01	22.823	35.566

Definición de X (predictoras) e y (a predecir)

```
In [19]: X_train = lag_df.iloc[:, :-1]
y_train = lag_df.iloc[:, -1]
```

Validación cruzada temporal

```
In [20]: train3, validation3 = ts_split(lag_df, train_size=0.75)
X_train3 = train3.iloc[:, :-1]
y_train3 = train3.iloc[:, -1]
X_val3 = validation3.iloc[:, :-1]
y_val3 = validation3.iloc[:, -1]

train2, validation2 = ts_split(train3, train_size=0.66)
X_train2 = train2.iloc[:, :-1]
y_train2 = train2.iloc[:, -1]
```

```

X_val2 = validation2.iloc[:, :-1]
y_val2 = validation2.iloc[:, -1]


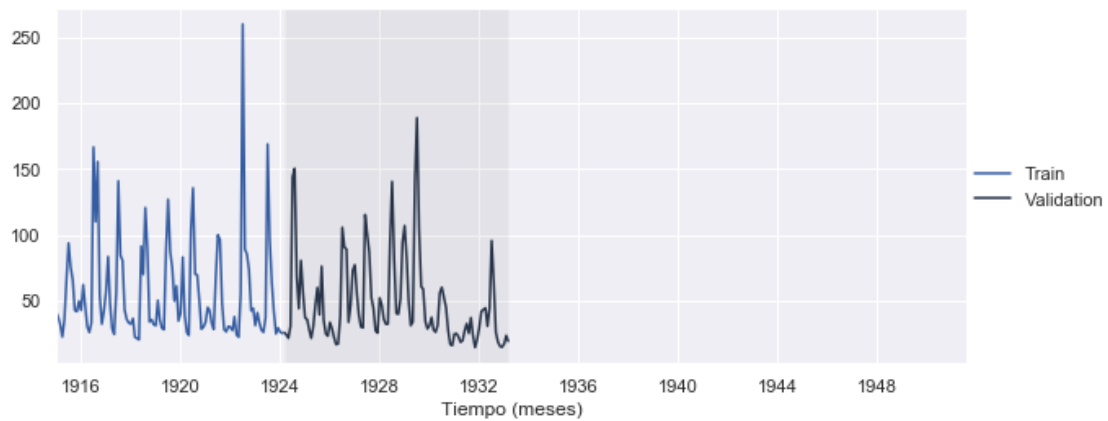
train1, validation1 = ts_split(train2, train_size=0.5)
X_train1 = train1.iloc[:, :-1]
y_train1 = train1.iloc[:, -1]
X_val1 = validation1.iloc[:, :-1]
y_val1 = validation1.iloc[:, -1]


```

```

In [21]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
plt.show()
tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
plt.show()
tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
plt.show()

```





K-Nearest Neighbors

```
In [22]: neighb = list(np.arange(2, 41, 1))
```

```
In [23]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN: k = %s | CVMSE = %.3f' % (best_knn['n_neighbors'],
                                           best_score))
```

```
Best KNN: k = 12 | CVMSE = 30.586
```

```
In [24]: knn_hyp = {"k": 12}
```

```

In [25]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid -
              knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)

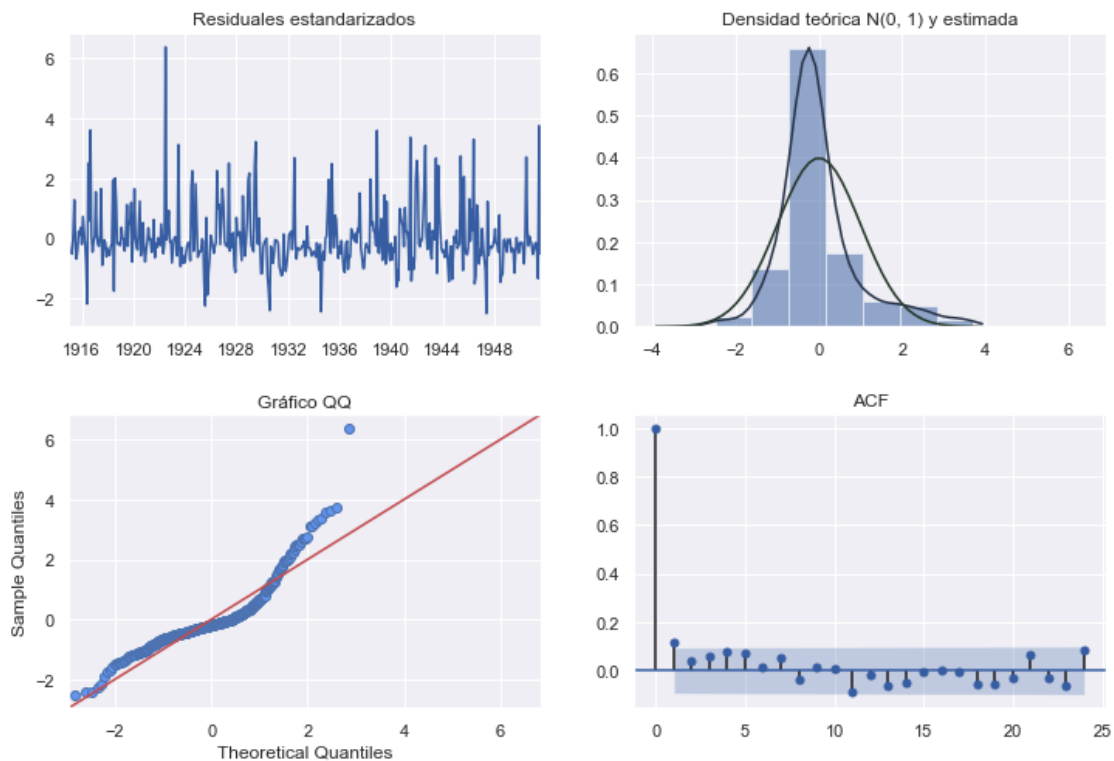
knn_forecast = iter_forecast(knn, y_train, 12, test)

```

```

In [26]: resid_diag(knn_resid)
plt.show()

```



```

In [27]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))

```

```

print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, knn_forecast))))
print("MAE (test): " + str(mae(test, knn_forecast)))
print("sMAPE (test): " + str(smape(test, knn_forecast)))

```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 2.903200667052171e-168

Test de Ljung-Box para k = 6 (p-valor): 0.03932935347286502

Test de Ljung-Box para k = 12 (p-valor): 0.07230606434407635

RMSE (test): 25.57542279856696


MAE (test): 16.631637648809523

sMAPE (test): 35.39380649389102

```

In [28]: forecast_plot(test, knn_forecast)
plt.show()

```



Regression tree

```

In [29]: depths = list(np.arange(1, 21, 1))

```



```

In [30]: best_score, best_tree = float("inf"), None

        for depth in depths:
            tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

            tree_model1 = tree.fit(X_train1, y_train1)
            tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))

            tree_model2 = tree.fit(X_val1, y_val1)
            tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
            rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

            tree_model3 = tree.fit(X_val2, y_val2)
            tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
            rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))

            cvrmse = mean([rmse1, rmse2, rmse3])
            if cvrmse < best_score:
                best_score, best_tree = cvrmse, tree.get_params()

        print('Best Tree: depth: %s | CV RMSE = %.3f' % (best_tree['max_depth'],
                                                         best_score))

```

Best Tree: depth: 1 | CV RMSE = 34.312

```

In [31]: tree_hyp = {"depth": 1}

```

```

In [32]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"], random_state=seed)
        tree_model = tree.fit(X_train, y_train)
        tree_fitted = pd.Series(tree_model.predict(X_train))
        tree_fitted
        tree_fitted.index = X_train.index
        tree_resid = y_train - tree_fitted
        tree_resid = (tree_resid -
                      tree_resid.mean()) / tree_resid.std()

        tree_jb_test = jb(tree_resid)
        tree_lb_test = lb(tree_resid)

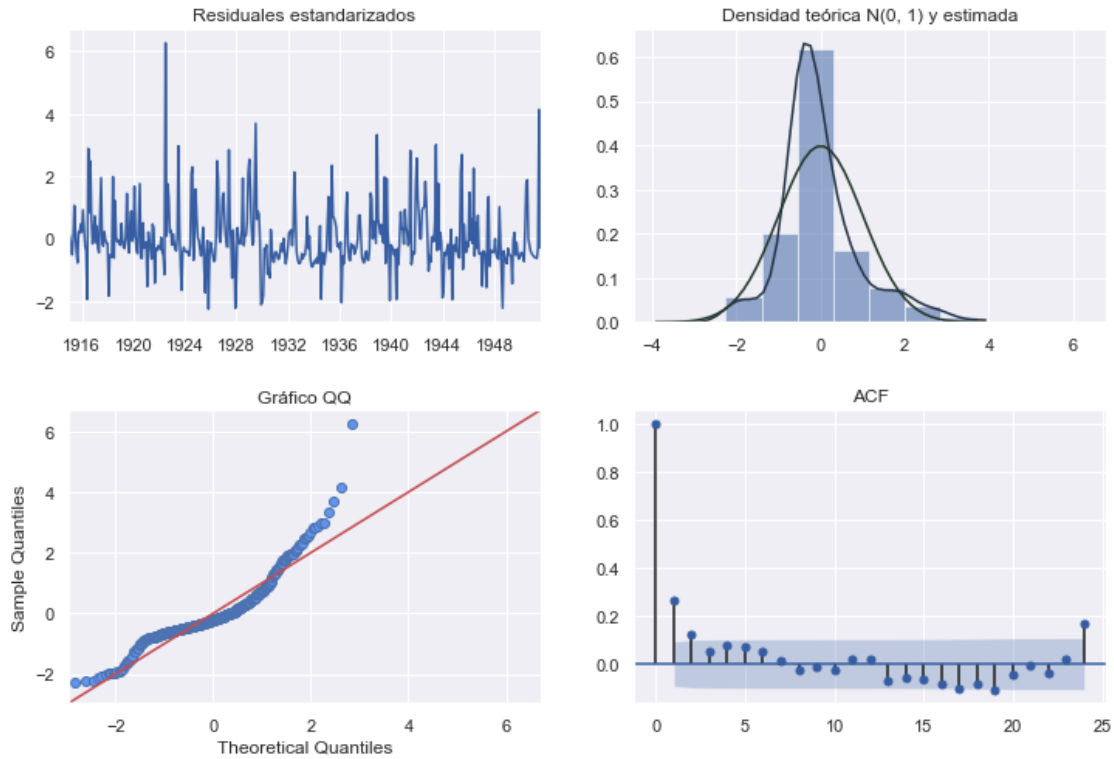
        tree_forecast = iter_forecast(tree, y_train, 12, test)

```

```

In [33]: resid_diag(tree_resid)
        plt.show()

```



```
In [34]: print("Decision Tree")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(tree_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(tree_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, tree_forecast))))
print("MAE (test): " + str(mae(test, tree_forecast)))
print("sMAPE (test): " + str(smape(test, tree_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 2.3423838676505707e-129

Test de Ljung-Box para k = 6 (p-valor): 2.0285871463335237e-07

Test de Ljung-Box para k = 12 (p-valor): 8.268260440498782e-06

```
RMSE (test): 25.83950361433136
MAE (test): 18.321662961679994
sMAPE (test): 40.46374917218229
```

```
In [35]: forecast_plot(test, tree_forecast)
plt.show()
```



Random Forest

```
In [37]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                       max_depth=depth,
                                       max_features=features,
                                       random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))
```

```

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features : %s | CV RMSE = %.3f' %
      (best_rf['n_estimators'], best_rf['max_depth'],
       best_rf['max_features'], best_score))

Best RF: n_estimators: 50 max_depth: 1.0 max_features : 0.8 | CV RMSE = 32.379

In [38]: rf_hyp = {"max_depth": 1,
                  "n_estimators": 50,
                  "max_features": 0.8}

In [39]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"], n_estimators=rf_hyp["n_estimators"],
                                   max_features=rf_hyp["max_features"], random_state=seed)

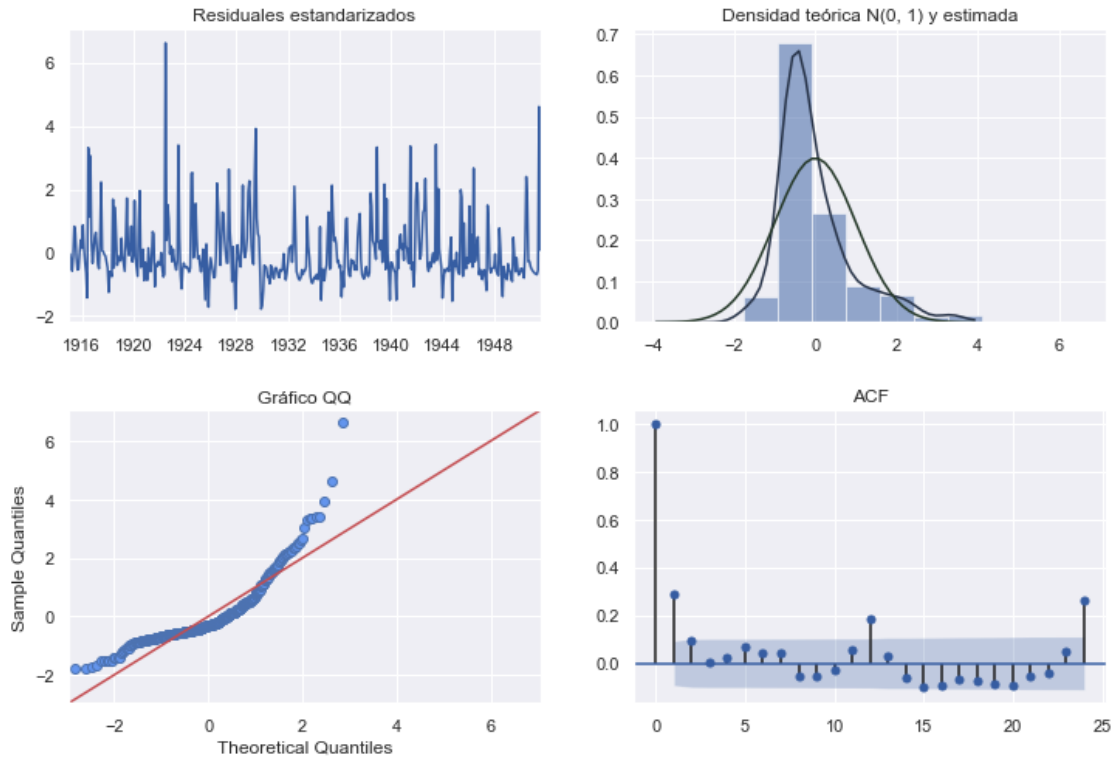
rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

In [40]: resid_diag(rf_resid)
plt.show()

```



```
In [41]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_forecast))))
print("MAE (test): " + str(mae(test, rf_forecast)))
print("sMAPE (test): " + str(smape(test, rf_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 1.7713368672108015e-226

Test de Ljung-Box para k = 6 (p-valor): 1.3674098923847824e-07

Test de Ljung-Box para k = 12 (p-valor): 4.580087633609218e-09

```
RMSE (test): 27.274163651871557
MAE (test): 18.873514748356094
sMAPE (test): 42.28556263637247
```

```
In [42]: forecast_plot(test, rf_forecast)
plt.show()
```



ElasticNet

```
In [43]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:, -1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
        rmse2 = sqrt(mse(validation2.iloc[:, -1], en_forecast))

        en_model3 = en.fit(X_val2, y_val2)
```

```

en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], en_forecast))

cvrmse = mean([rmse1, rmse2, rmse3])
if cvrmse < best_score:
    best_score, best_en = cvrmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

Best EN: ratio: 0.1 alpha: 100 | CV RMSE = 31.974

In [44]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 100}

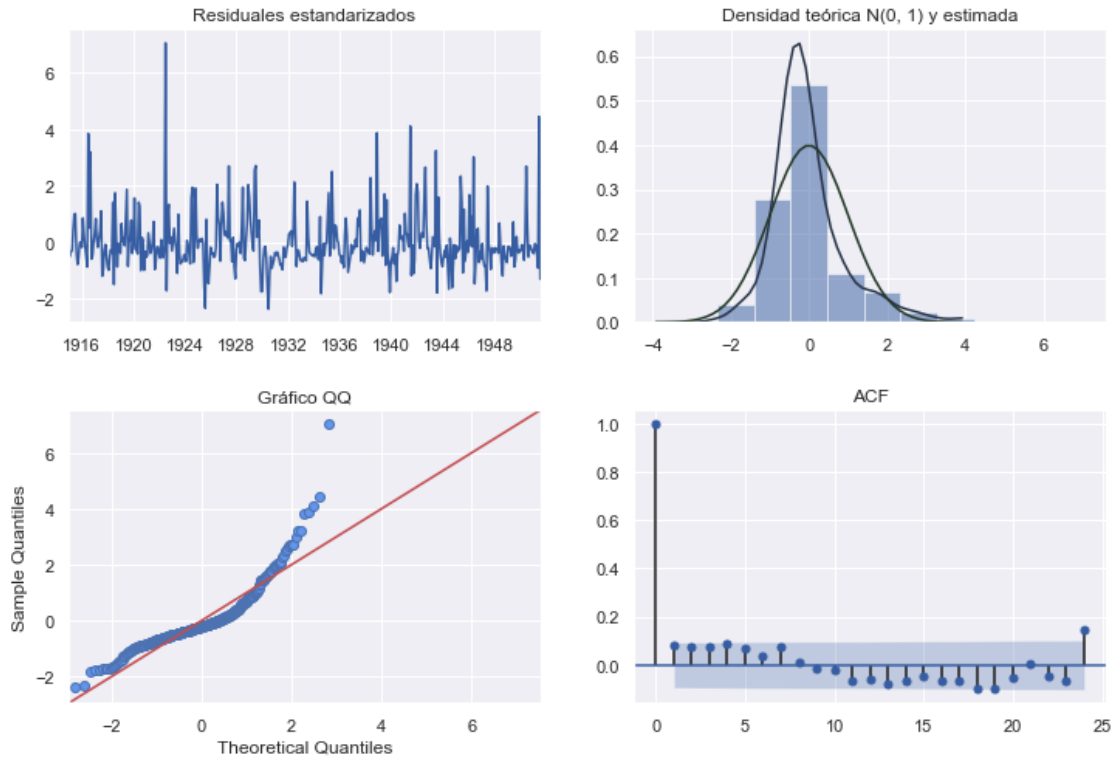
In [45]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

In [46]: resid_diag(en_resid)
plt.show()

```



```
In [47]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_forecast))))
print("MAE (test): " + str(mae(test, en_forecast)))
print("sMAPE (test): " + str(smape(test, en_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 3.034582288460461e-305

Test de Ljung-Box para k = 6 (p-valor): 0.018580289030905108

Test de Ljung-Box para k = 12 (p-valor): 0.0377027495242834


```


RMSE (test): 30.11545109033476
MAE (test): 22.0843920027641
sMAPE (test): 48.59639886725879

```

```

In [48]: forecast_plot(test, en_forecast)
plt.show()

```



Serie temporal transformada con Box-Cox como estructura de DataFrame

```

In [49]: lag_df = series_to_supervised(bc_train, n_in=12)
lag_df.head(5)

```

```

Out[49]:
          var1(t-12)  var1(t-11)  var1(t-10)  var1(t-9)  var1(t-8)  \
index
1915-01-01    2.056547    2.034131    1.997200    1.876680    1.825973
1915-02-01    2.034131    1.997200    1.876680    1.825973    1.953956
1915-03-01    1.997200    1.876680    1.825973    1.953956    2.093966
1915-04-01    1.876680    1.825973    1.953956    2.093966    2.140314
1915-05-01    1.825973    1.953956    2.093966    2.140314    2.193369

          var1(t-7)  var1(t-6)  var1(t-5)  var1(t-4)  var1(t-3)  var1(t-2)  \
index
1915-01-01    1.953956    2.093966    2.140314    2.193369    2.076357    1.959344
1915-02-01    2.093966    2.140314    2.193369    2.076357    1.959344    2.041142
1915-03-01    2.140314    2.193369    2.076357    1.959344    2.041142    1.992191
1915-04-01    2.193369    2.076357    1.959344    2.041142    1.992191    1.965730
1915-05-01    2.076357    1.959344    2.041142    1.992191    1.965730    1.924627

```

	var1(t-1)	var1(t)
index		
1915-01-01	2.041142	1.992191
1915-02-01	1.992191	1.965730
1915-03-01	1.965730	1.924627
1915-04-01	1.924627	1.828236
1915-05-01	1.828236	1.952324

Definición de X (predictoras) e y (a predecir)

```
In [50]: X_train = lag_df.iloc[:, :-1]
        y_train = lag_df.iloc[:, -1]
```


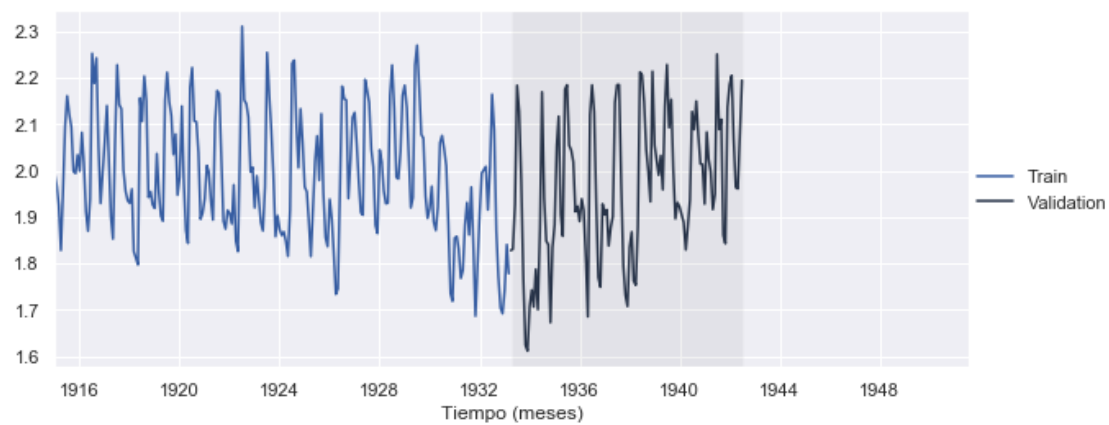

Validación cruzada temporal

```
In [51]: train3, validation3 = ts_split(lag_df, train_size=0.75)
        X_train3 = train3.iloc[:, :-1]
        y_train3 = train3.iloc[:, -1]
        X_val3 = validation3.iloc[:, :-1]
        y_val3 = validation3.iloc[:, -1]

        train2, validation2 = ts_split(train3, train_size=0.66)
        X_train2 = train2.iloc[:, :-1]
        y_train2 = train2.iloc[:, -1]
        X_val2 = validation2.iloc[:, :-1]
        y_val2 = validation2.iloc[:, -1]

        train1, validation1 = ts_split(train2, train_size=0.5)
        X_train1 = train1.iloc[:, :-1]
        y_train1 = train1.iloc[:, -1]
        X_val1 = validation1.iloc[:, :-1]
        y_val1 = validation1.iloc[:, -1]

In [52]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
        plt.show()
        tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
        plt.show()
        tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
        plt.show()
```



K-Nearest Neighbors

```
In [53]: neighb = list(np.arange(2, 41, 1))
```

```
In [54]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVRMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

Best KNN 10 CVRMSE = 0.136

```
In [55]: knn_hyp = {"k": 10}
```

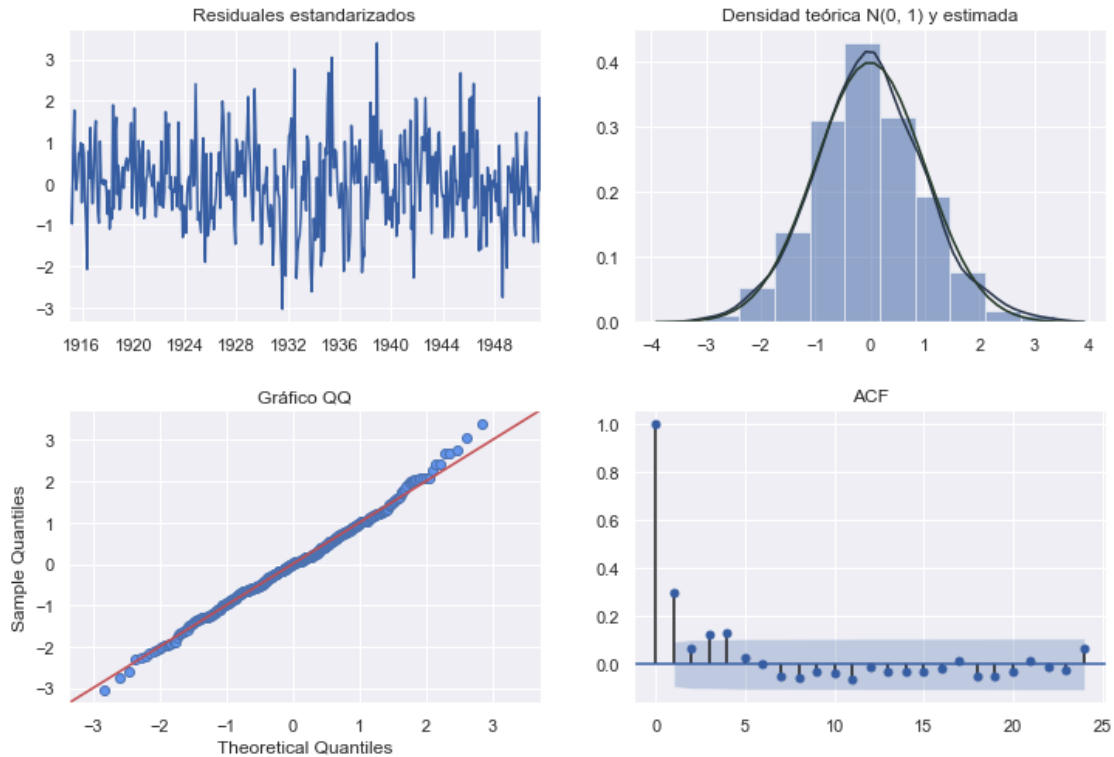
```
In [56]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)
```

```
knn_forecast = iter_forecast(knn, y_train, 12, test)
```

```
In [57]: knn_box_forecast = (bc_param * knn_forecast + 1) ** (1 / bc_param)
```

```
In [58]: resid_diag(knn_resid)
plt.show()
```



```
In [59]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, knn_box_forecast))))
print("MAE (test): " + str(mae(test, knn_box_forecast)))
print("sMAPE (test): " + str(smape(test, knn_box_forecast)))
```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 0.14929158900292663

Test de Ljung-Box para k = 6 (p-valor): 6.173994051775582e-10

Test de Ljung-Box para k = 12 (p-valor): 2.2780049116050887e-08

```
RMSE (test): 23.826827407781153
MAE (test): 16.289851321109712
sMAPE (test): 35.04950163941397
```

```
In [60]: forecast_plot(test, knn_box_forecast)
plt.show()
```



Regression tree

```
In [61]: depths = list(np.arange(1, 21, 1))

In [62]: best_score, best_tree = float("inf"), None

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)
    tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))

    tree_model2 = tree.fit(X_val1, y_val1)
    tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

    tree_model3 = tree.fit(X_val2, y_val2)
    tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))
```

```

        cvrmse = mean([rmse1, rmse2, rmse3])
        if cvrmse < best_score:
            best_score, best_tree = cvrmse, tree.get_params()

    print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

Best Tree 5 CV RMSE = 0.152

In [63]: tree_hyp = {"depth": 5}

In [64]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                     random_state=seed)
        tree_model = tree.fit(X_train, y_train)
        tree_fitted = pd.Series(tree_model.predict(X_train))
        tree_fitted
        tree_fitted.index = X_train.index
        tree_resid = y_train - tree_fitted
        tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

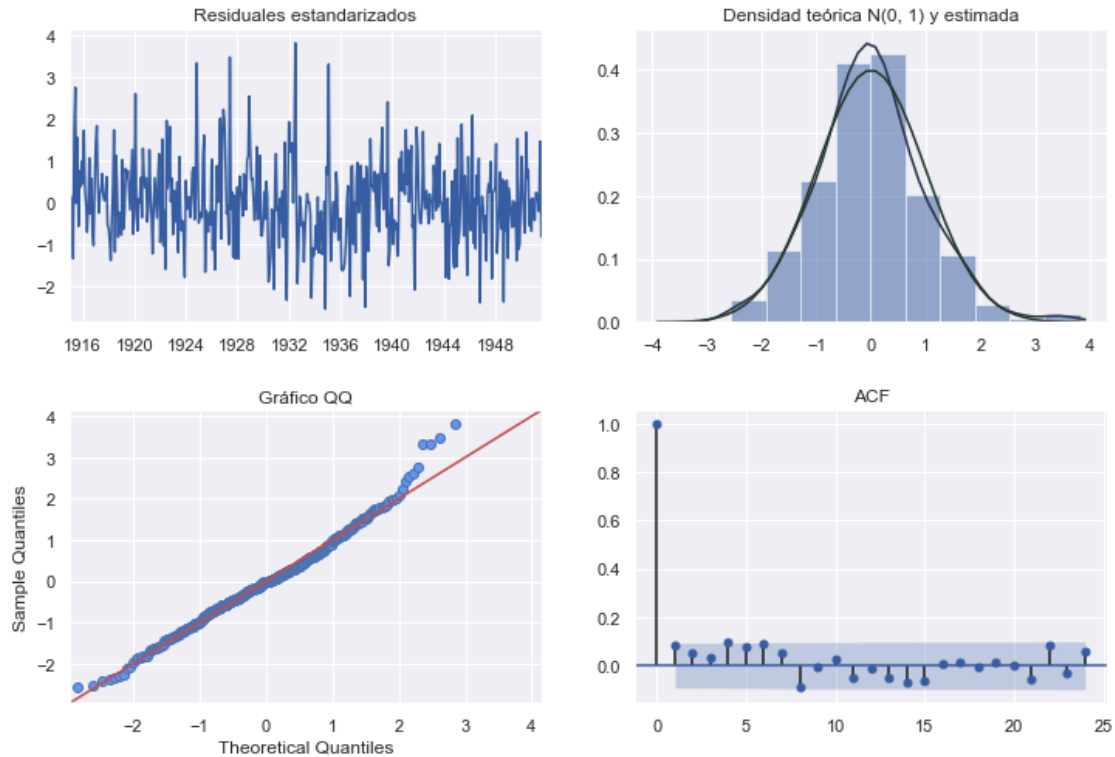
        tree_jb_test = jb(tree_resid)
        tree_lb_test = lb(tree_resid)

        tree_forecast = iter_forecast(tree, y_train, 12, test)

In [65]: tree_box_forecast = (bc_param * tree_forecast + 1) ** (1 / bc_param)

In [66]: resid_diag(tree_resid)
        plt.show()

```



```
In [67]: print("Decision Tree")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(tree_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(tree_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, tree_box_forecast))))
print("MAE (test): " + str(mae(test, tree_box_forecast)))
print("sMAPE (test): " + str(smape(test, tree_box_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 1.566066446696135e-05

Test de Ljung-Box para k = 6 (p-valor): 0.020923869379414626

Test de Ljung-Box para k = 12 (p-valor): 0.0421623493635886


```
RMSE (test): 24.744712241704015
MAE (test): 16.56474204351321
sMAPE (test): 36.708903576558605
```

```
In [68]: forecast_plot(test, tree_box_forecast)
plt.show()
```



Random Forest

```
In [69]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                      max_depth=depth,
                                      max_features=features,
                                      random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))
```

```

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f' %
      (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_score))

Best RF: n_estimators: 50 max_depth: 7.315789473684211 max_features :0.8 | CV RMSE = 0.140

In [70]: rf_hyp = {"max_depth": 7.32,
                  "n_estimators": 50,
                  "max_features": 0.8}

In [71]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                   n_estimators=rf_hyp["n_estimators"],
                                   max_features=rf_hyp["max_features"],
                                   random_state=seed)

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

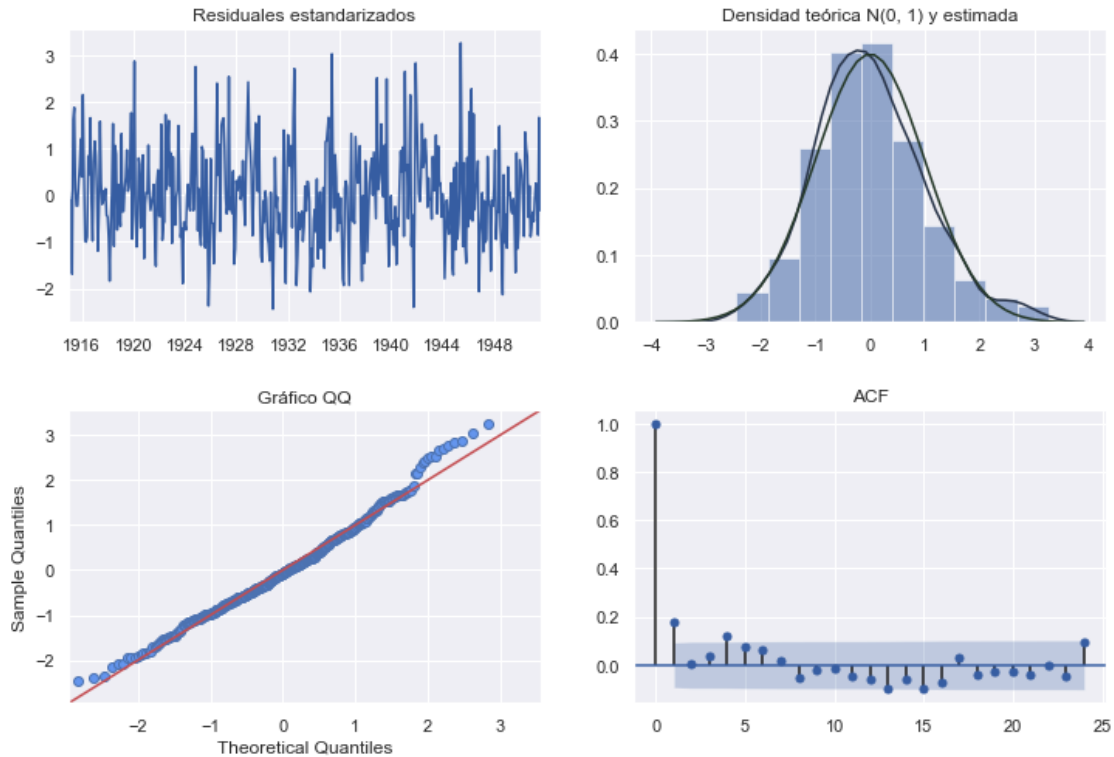
rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

In [72]: rf_box_forecast = (bc_param * rf_forecast + 1) ** (1 / bc_param)

In [73]: resid_diag(rf_resid)
plt.show()

```



```
In [74]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_box_forecast))))
print("MAE (test): " + str(mae(test, rf_box_forecast)))
print("sMAPE (test): " + str(smape(test, rf_box_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 0.0009343986727170103

Test de Ljung-Box para k = 6 (p-valor): 0.0007964689182289883

Test de Ljung-Box para k = 12 (p-valor): 0.0016563699647625693

```
RMSE (test): 30.496564708279912
MAE (test): 17.28236557049287
sMAPE (test): 37.02239483180497
```

```
In [75]: forecast_plot(test, rf_box_forecast)
plt.show()
```



ElasticNet

```
In [76]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:, -1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
        rmse2 = sqrt(mse(validation2.iloc[:, -1], en_forecast))

        en_model3 = en.fit(X_val2, y_val2)
```

```

en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], en_forecast))

cvrmse = mean([rmse1, rmse2, rmse3])
if cvrmse < best_score:
    best_score, best_en = cvrmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

Best EN: ratio: 0.1 alpha: 0 | CV RMSE = 0.140

In [77]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0}

In [78]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

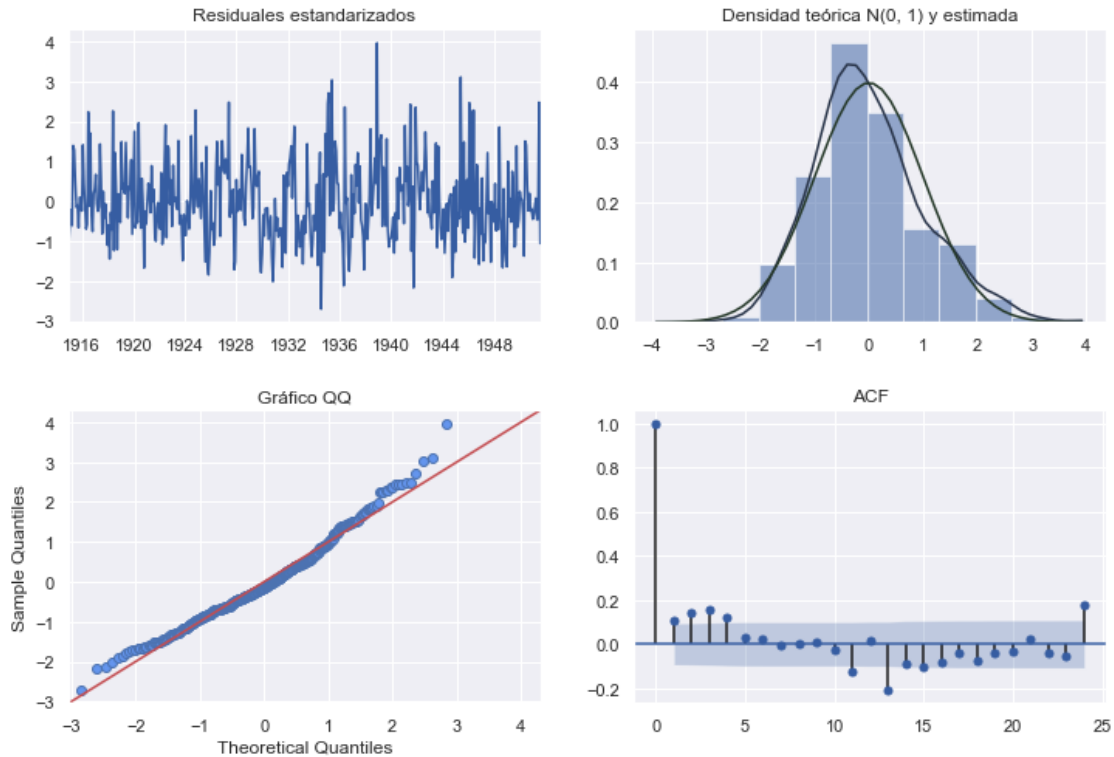
en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

In [79]: en_box_forecast = (bc_param * en_forecast + 1) ** (1 / bc_param)

In [80]: resid_diag(en_resid)
plt.show()

```



```
In [81]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_box_forecast))))
print("MAE (test): " + str(mae(test, en_box_forecast)))
print("sMAPE (test): " + str(smape(test, en_box_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 6.76771127119615e-07

Test de Ljung-Box para k = 6 (p-valor): 4.748066502114843e-05

Test de Ljung-Box para k = 12 (p-valor): 1.0413255691897458e-07

```
RMSE (test): 33.682465160124686
MAE (test): 20.119056207537295
sMAPE (test): 42.9938082467703
```

```
In [82]: forecast_plot(test, en_box_forecast)
plt.show()
```



Serie temporal desestacionalizada como estructura de DataFrame

```
In [83]: lag_df = series_to_supervised(deseas, n_in=12)
lag_df.head(5)
```

```
Out[83]:
```

	var1(t-12)	var1(t-11)	var1(t-10)	var1(t-9)	var1(t-8)	\
index						
1915-01-01	68.056741	55.487990	59.745262	49.642875	47.097132	
1915-02-01	55.487990	59.745262	49.642875	47.097132	19.921335	
1915-03-01	59.745262	49.642875	47.097132	19.921335	2.252072	
1915-04-01	49.642875	47.097132	19.921335	2.252072	56.538266	
1915-05-01	47.097132	19.921335	2.252072	56.538266	96.012723	

	var1(t-7)	var1(t-6)	var1(t-5)	var1(t-4)	var1(t-3)	var1(t-2)	\
index							
1915-01-01	19.921335	2.252072	56.538266	96.012723	70.726177	55.973661	
1915-02-01	2.252072	56.538266	96.012723	70.726177	55.973661	64.087769	
1915-03-01	56.538266	96.012723	70.726177	55.973661	64.087769	54.520741	
1915-04-01	96.012723	70.726177	55.973661	64.087769	54.520741	43.028990	
1915-05-01	70.726177	55.973661	64.087769	54.520741	43.028990	49.126262	

	var1(t-1)	var1(t)
index		
1915-01-01	64.087769	54.520741
1915-02-01	54.520741	43.028990
1915-03-01	43.028990	49.126262
1915-04-01	49.126262	45.564875
1915-05-01	45.564875	60.009132

Definición de X (predictoras) e y (a predecir)

```
In [84]: X_train = lag_df.iloc[:, :-1]
        y_train = lag_df.iloc[:, -1]
```

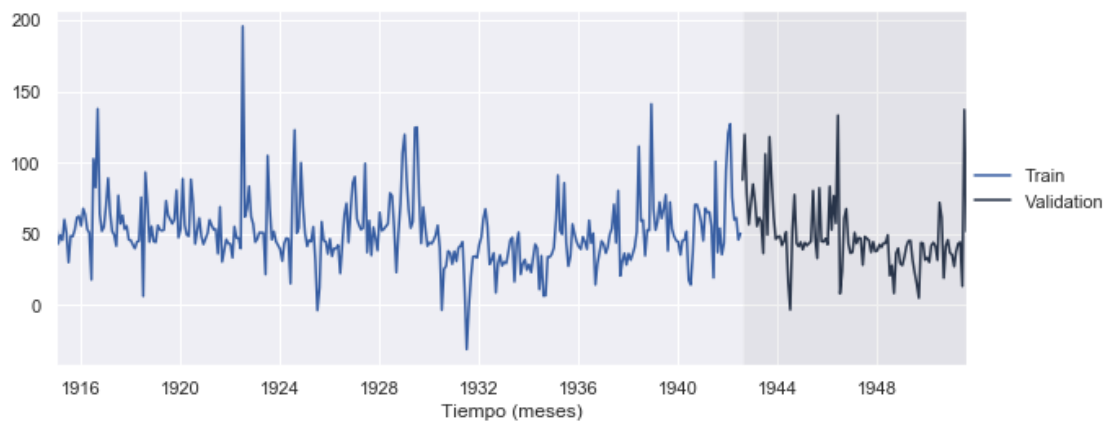

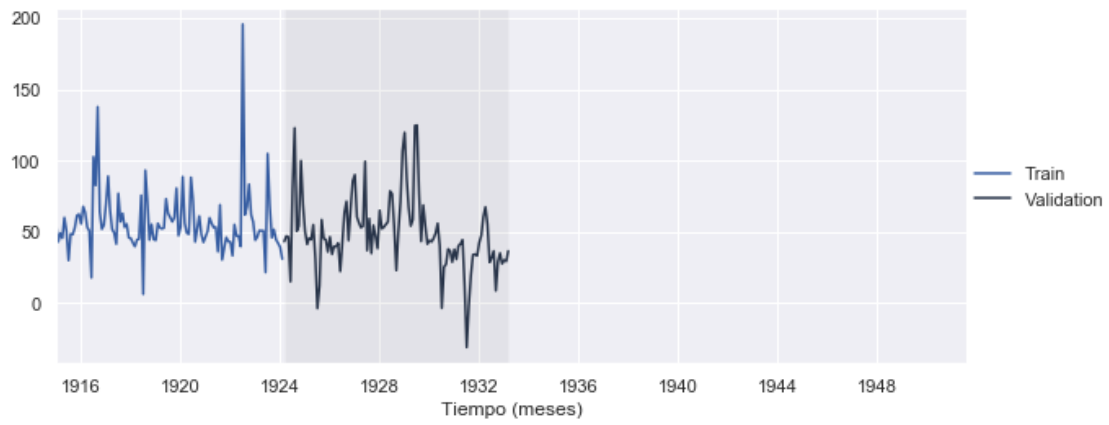
Validación cruzada temporal

```
In [85]: train3, validation3 = ts_split(lag_df, train_size=0.75)
        X_train3 = train3.iloc[:, :-1]
        y_train3 = train3.iloc[:, -1]
        X_val3 = validation3.iloc[:, :-1]
        y_val3 = validation3.iloc[:, -1]

        train2, validation2 = ts_split(train3, train_size=0.66)
        X_train2 = train2.iloc[:, :-1]
        y_train2 = train2.iloc[:, -1]
        X_val2 = validation2.iloc[:, :-1]
        y_val2 = validation2.iloc[:, -1]

        train1, validation1 = ts_split(train2, train_size=0.5)
        X_train1 = train1.iloc[:, :-1]
        y_train1 = train1.iloc[:, -1]
        X_val1 = validation1.iloc[:, :-1]
        y_val1 = validation1.iloc[:, -1]

In [86]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
        plt.show()
        tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
        plt.show()
        tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
        plt.show()
```

K-Nearest Neighbors

```
In [87]: neighb = list(np.arange(2, 41, 1))
```

```
In [88]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVRMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

```
Best KNN 10 CVRMSE = 24.538
```

```
In [ ]: knn_hyp = {"k": 10}
```

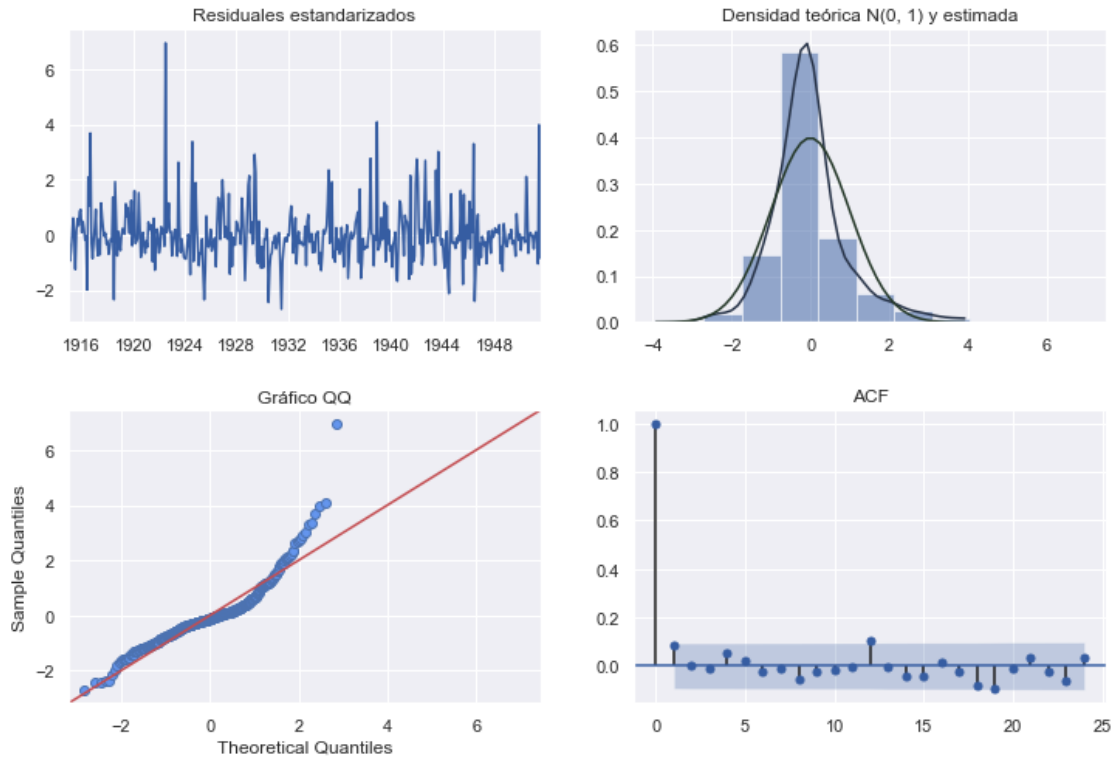
```
In [89]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
    knn_model = knn.fit(X_train, y_train)
    knn_fitted = pd.Series(knn_model.predict(X_train))
    knn_fitted
    knn_fitted.index = X_train.index
    knn_resid = y_train - knn_fitted
    knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

    knn_jb_test = jb(knn_resid)
    knn_lb_test = lb(knn_resid)

    knn_forecast = iter_forecast(knn, y_train, 12, test)

In [90]: knn_seas_forecast = knn_forecast + seasonality_forecast

In [91]: resid_diag(knn_resid)
    plt.show()
```



```
In [92]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, knn_seas_forecast))))
print("MAE (test): " + str(mae(test, knn_seas_forecast)))
print("sMAPE (test): " + str(smape(test, knn_seas_forecast)))
```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 3.221140212318335e-238

Test de Ljung-Box para k = 6 (p-valor): 0.6861542427252421

Test de Ljung-Box para k = 12 (p-valor): 0.5852275392995411

```
RMSE (test): 35.73070318812319
MAE (test): 23.860184319978462
sMAPE (test): 52.04783996158255
```

```
In [93]: forecast_plot(test, knn_seas_forecast)
plt.show()
```



Regression tree

```
In [94]: depths = list(np.arange(1, 21, 1))
```

```
In [95]: best_score, best_tree = float("inf"), None
```

```
for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)
    tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))

    tree_model2 = tree.fit(X_val1, y_val1)
    tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

    tree_model3 = tree.fit(X_val2, y_val2)
    tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))
```

```

        cvrmse = mean([rmse1, rmse2, rmse3])

        if cvrmse < best_score:
            best_score, best_tree = cvrmse, tree.get_params()

    print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

Best Tree 3 CV RMSE = 24.941

In [96]: tree_hyp = {"depth": 3}

In [97]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                     random_state=seed)
        tree_model = tree.fit(X_train, y_train)
        tree_fitted = pd.Series(tree_model.predict(X_train))
        tree_fitted
        tree_fitted.index = X_train.index
        tree_resid = y_train - tree_fitted
        tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

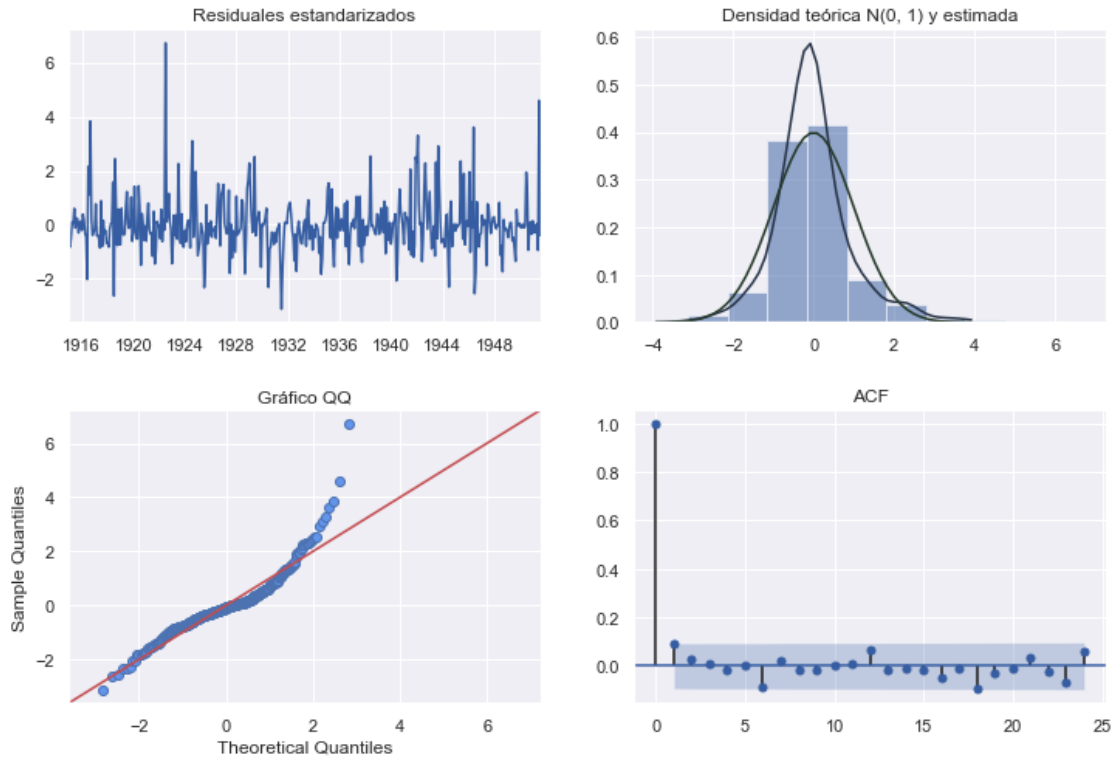
        tree_jb_test = jb(tree_resid)
        tree_lb_test = lb(tree_resid)

        tree_forecast = iter_forecast(tree, y_train, 12, test)

In [98]: tree_seas_forecast = tree_forecast + seasonality_forecast

In [99]: resid_diag(tree_resid)
        plt.show()

```



```
In [100]: print("Decision Tree")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(tree_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(tree_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, tree_seas_forecast))))
          print("MAE (test): " + str(mae(test, tree_seas_forecast)))
          print("sMAPE (test): " + str(smape(test, tree_seas_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 2.3249335220687302e-200

Test de Ljung-Box para k = 6 (p-valor): 0.32816725368366206

Test de Ljung-Box para k = 12 (p-valor): 0.648859080395447

```
RMSE (test): 35.32047965264542
MAE (test): 23.998700669257506
sMAPE (test): 54.28693445110659
```

```
In [101]: forecast_plot(test, tree_seas_forecast)
plt.show()
```



Random Forest

```
In [102]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                       max_depth=depth,
                                       max_features=features,
                                       random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))
```

```

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f'
      % (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_

Best RF: n_estimators: 50 max_depth: 1.0 max_features :1 | CV RMSE = 24.151

In [103]: rf_hyp = {"max_depth": 1,
                    "n_estimators": 50,
                    "max_features": 1}

In [104]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                     n_estimators=rf_hyp["n_estimators"],
                                     max_features=rf_hyp["max_features"])

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

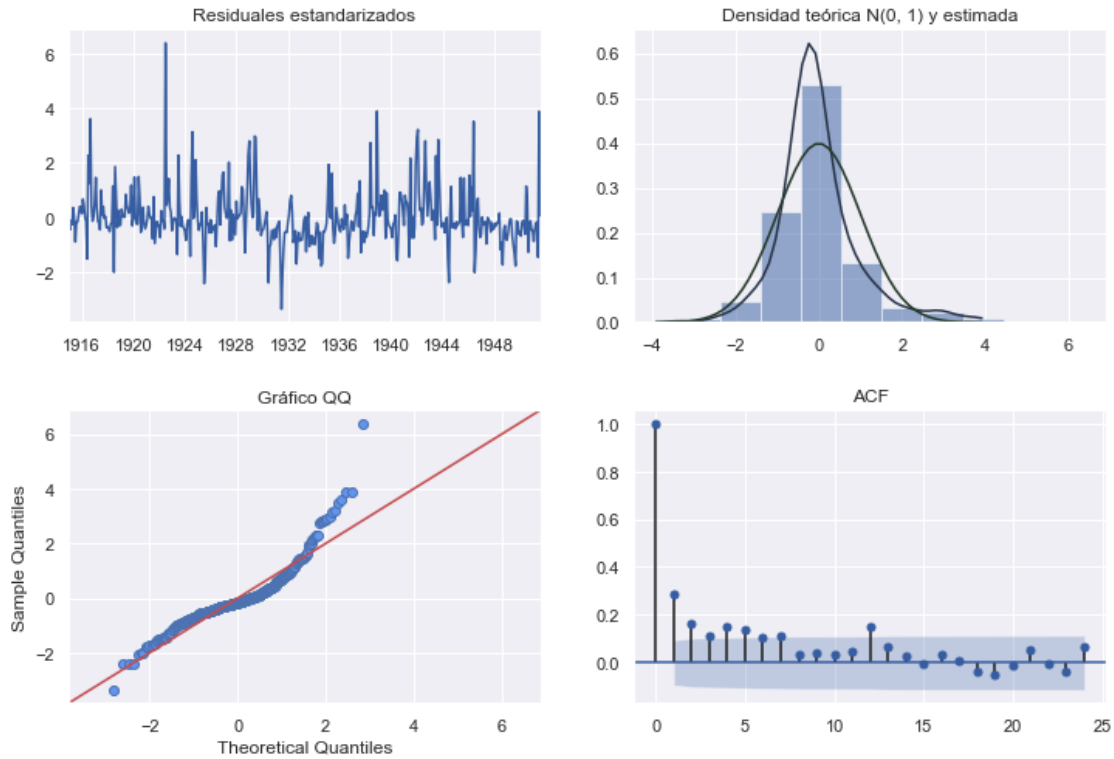
rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

In [105]: rf_seas_forecast = rf_forecast + seasonality_forecast

In [106]: resid_diag(rf_resid)
plt.show()

```

```
In [107]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_seas_forecast))))
print("MAE (test): " + str(mae(test, rf_seas_forecast)))
print("sMAPE (test): " + str(smape(test, rf_seas_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 1.3404599347819535e-167

Test de Ljung-Box para k = 6 (p-valor): 3.8209202885056384e-15

Test de Ljung-Box para k = 12 (p-valor): 3.08333506870211e-15

```
RMSE (test): 37.221632138684946
MAE (test): 26.05070483344991
sMAPE (test): 50.40460863038776
```

```
In [108]: forecast_plot(test, rf_seas_forecast)
plt.show()
```



ElasticNet

```
In [109]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:,-1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
        rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

        en_model3 = en.fit(X_val2, y_val2)
```

```

en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], en_forecast))

cvrmse = mean([rmse1, rmse2, rmse3])
if cvrmse < best_score:
    best_score, best_en = cvrmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

```

Best EN: ratio: 0.7 alpha: 50 | CV RMSE = 24.458

```

In [110]: en_hyp = {"l1_ratio": 0.7,
                  "alpha": 50}

```

```

In [111]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

```

```

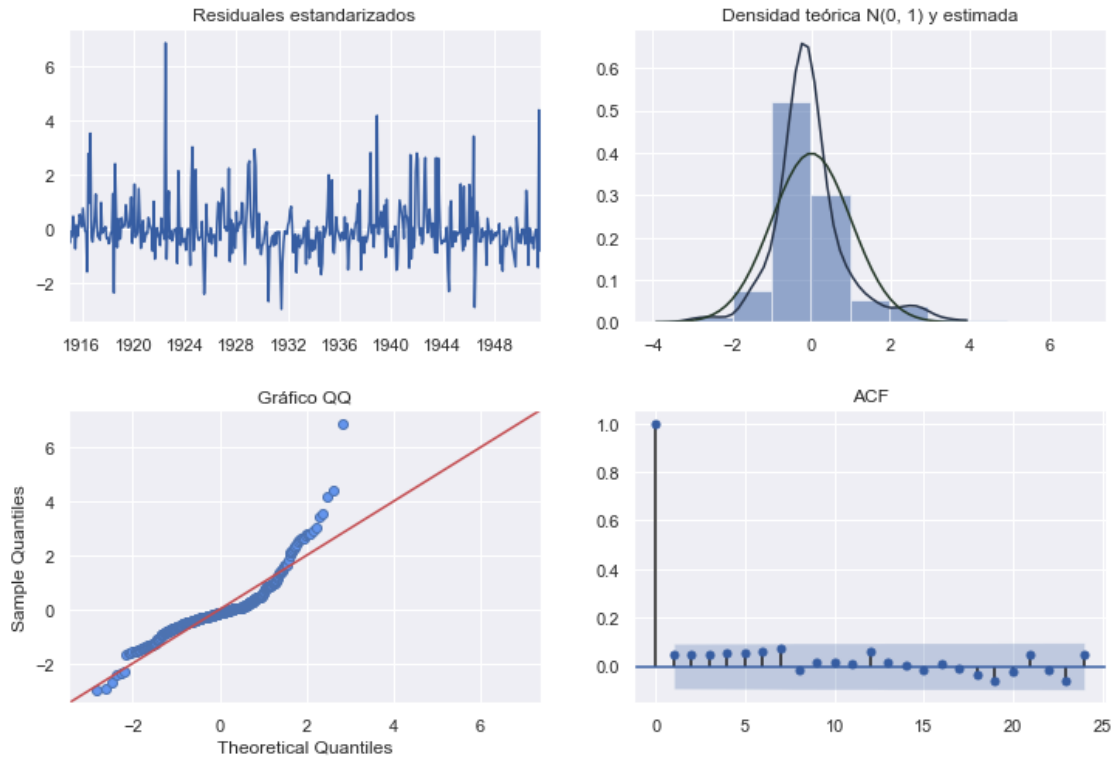
In [112]: en_seas_forecast = en_forecast + seasonality_forecast

```

```

In [113]: resid_diag(en_resid)
plt.show()

```



```
In [114]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_seas_forecast))))
print("MAE (test): " + str(mae(test, en_seas_forecast)))
print("sMAPE (test): " + str(smape(test, en_seas_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 1.8221992633454572e-265

Test de Ljung-Box para k = 6 (p-valor): 0.2116590961603848

Test de Ljung-Box para k = 12 (p-valor): 0.5484740779269146

```


RMSE (test): 35.879726444295535
MAE (test): 24.27506662692949
sMAPE (test): 47.12662561864641

```

```

In [115]: forecast_plot(test, en_seas_forecast)
plt.show()

```



Serie temporal diferenciada como estructura de DataFrame

```

In [116]: lag_df = series_to_supervised(diff_train, n_in=12)
lag_df.head(5)

```

```

Out[116]:
      var1(t-12)  var1(t-11)  var1(t-10)  var1(t-9)  var1(t-8)  \
index
1915-02-01    -5.324    -7.334    -15.716    -4.247    13.139
1915-03-01    -7.334   -15.716     -4.247    13.139    30.270
1915-04-01   -15.716     -4.247    13.139    30.270    17.840
1915-05-01     -4.247    13.139    30.270    17.840    29.733
1915-06-01    13.139    30.270    17.840    29.733   -52.981

      var1(t-7)  var1(t-6)  var1(t-5)  var1(t-4)  var1(t-3)  var1(t-2)  \
index
1915-02-01    30.270    17.840    29.733   -52.981   -24.098    14.980
1915-03-01    17.840    29.733   -52.981   -24.098    14.980   -9.798
1915-04-01    29.733   -52.981   -24.098    14.980   -9.798   -4.247
1915-05-01   -52.981   -24.098    14.980   -9.798   -4.247   -5.494
1915-06-01   -24.098    14.980   -9.798   -4.247   -5.494   -9.175

```

	var1(t-1)	var1(t)
index		
1915-02-01	-9.798	-4.247
1915-03-01	-4.247	-5.494
1915-04-01	-5.494	-9.175
1915-05-01	-9.175	12.743
1915-06-01	12.743	30.724

Definición de X (predictoras) e y (a predecir)

```
In [117]: X_train = lag_df.iloc[:, :-1]
          y_train = lag_df.iloc[:, -1]
```




Validación cruzada temporal

```
In [118]: train3, validation3 = ts_split(lag_df, train_size=0.75)
          X_train3 = train3.iloc[:, :-1]
          y_train3 = train3.iloc[:, -1]
          X_val3 = validation3.iloc[:, :-1]
          y_val3 = validation3.iloc[:, -1]

          train2, validation2 = ts_split(train3, train_size=0.66)
          X_train2 = train2.iloc[:, :-1]
          y_train2 = train2.iloc[:, -1]
          X_val2 = validation2.iloc[:, :-1]
          y_val2 = validation2.iloc[:, -1]

          train1, validation1 = ts_split(train2, train_size=0.5)
          X_train1 = train1.iloc[:, :-1]
          y_train1 = train1.iloc[:, -1]
          X_val1 = validation1.iloc[:, :-1]
          y_val1 = validation1.iloc[:, -1]

In [119]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
          plt.show()
```



K-Nearest Neighbors

```
In [120]: neighb = list(np.arange(2, 41, 1))
```

```
In [121]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVRMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

Best KNN 8 CVRMSE = 33.097

```
In [123]: knn_hyp = {"k": 8}
```

```
In [124]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)

knn_forecast = iter_forecast(knn, y_train, 12, test)
```

```
In [125]: knn_undiff_forecast = []
last_obser = train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + knn_forecast[i]
```



```

knn_undiff_forecast.append(yhat)
last_obser = yhat

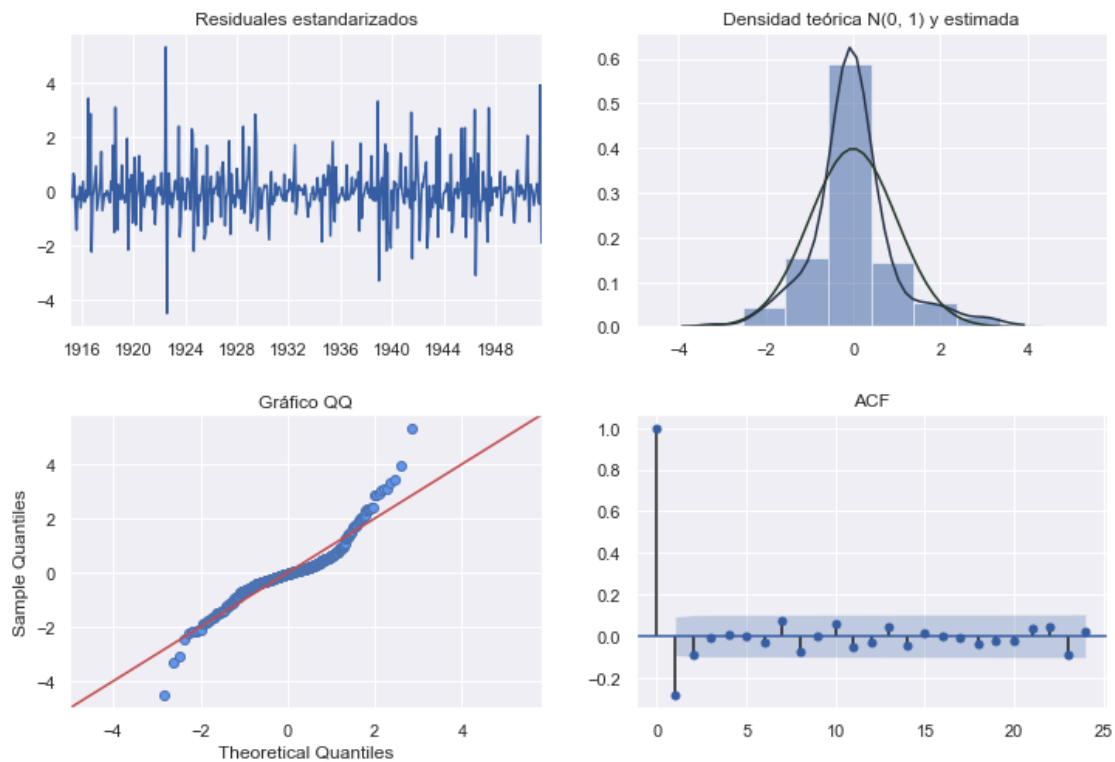
knn_undiff_forecast = pd.Series(knn_undiff_forecast)
knn_undiff_forecast.index = test.index

```

```

In [126]: resid_diag(knn_resid)
plt.show()

```



```

In [127]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, knn_undiff_forecast))))
print("MAE (test): " + str(mae(test, knn_undiff_forecast)))
print("sMAPE (test): " + str(smape(test, knn_undiff_forecast)))

```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 1.917901228467001e-79

Test de Ljung-Box para k = 6 (p-valor): 5.776945665723931e-07


Test de Ljung-Box para k = 12 (p-valor): 5.455386637939246e-06

RMSE (test): 177.22695016002845

MAE (test): 153.15659709821418

sMAPE (test): 121.110202320245

```
In [128]: forecast_plot(test, knn_undiff_forecast)
plt.show()
```



Regression tree

```
In [129]: depths = list(np.arange(1, 21, 1))
```

```
In [130]: best_score, best_tree = float("inf"), None
```

```
    for depth in depths:
        tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

        tree_model1 = tree.fit(X_train1, y_train1)
        tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))
```

```

tree_model2 = tree.fit(X_val1, y_val1)
tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

tree_model3 = tree.fit(X_val2, y_val2)
tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_tree = cvmse, tree.get_params()

print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

```

Best Tree 2 CV RMSE = 30.315

```
In [131]: tree_hyp = {"depth": 2}
```

```
In [132]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                     random_state=seed)
tree_model = tree.fit(X_train, y_train)
tree_fitted = pd.Series(tree_model.predict(X_train))
tree_fitted
tree_fitted.index = X_train.index
tree_resid = y_train - tree_fitted
tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

tree_jb_test = jb(tree_resid)
tree_lb_test = lb(tree_resid)

```

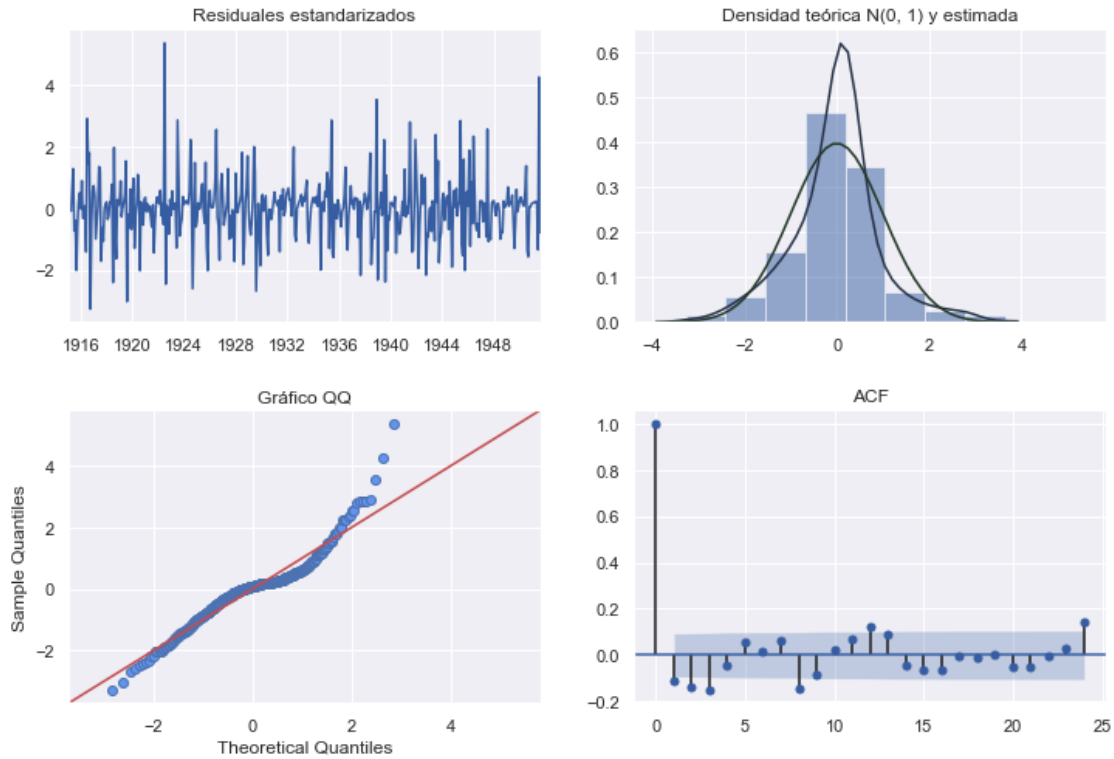
```
tree_forecast = iter_forecast(tree, y_train, 12, test)
```

```
In [133]: tree_undiff_forecast = []
last_obser = train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + tree_forecast[i]
    tree_undiff_forecast.append(yhat)
    last_obser = yhat

tree_undiff_forecast = pd.Series(tree_undiff_forecast)
tree_undiff_forecast.index = test.index

```

```
In [134]: resid_diag(tree_resid)
plt.show()
```



```
In [135]: print("Decision Tree")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(tree_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(tree_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, tree_undiff_forecast))))
          print("MAE (test): " + str(mae(test, tree_undiff_forecast)))
          print("sMAPE (test): " + str(smape(test, tree_undiff_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 1.2372416779099734e-51

Test de Ljung-Box para k = 6 (p-valor): 0.00023462568230369247

Test de Ljung-Box para k = 12 (p-valor): 6.963358532075138e-07

```
RMSE (test): 114.28259407187224
MAE (test): 93.38271343999499
sMAPE (test): 163.26939411832194
```

```
In [136]: forecast_plot(test, tree_undiff_forecast)
plt.show()
```



Random Forest

```
In [137]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                      max_depth=depth,
                                      max_features=features,
                                      random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))
```

```

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f'
      % (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_

Best RF: n_estimators: 200 max_depth: 19.94736842105263 max_features :0.8 | CV RMSE = 28.999

In [138]: rf_hyp = {"max_depth": 19.95,
                  "n_estimators": 200,
                  "max_features": 0.8}

In [139]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                   n_estimators=rf_hyp["n_estimators"],
                                   max_features=rf_hyp["max_features"])

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

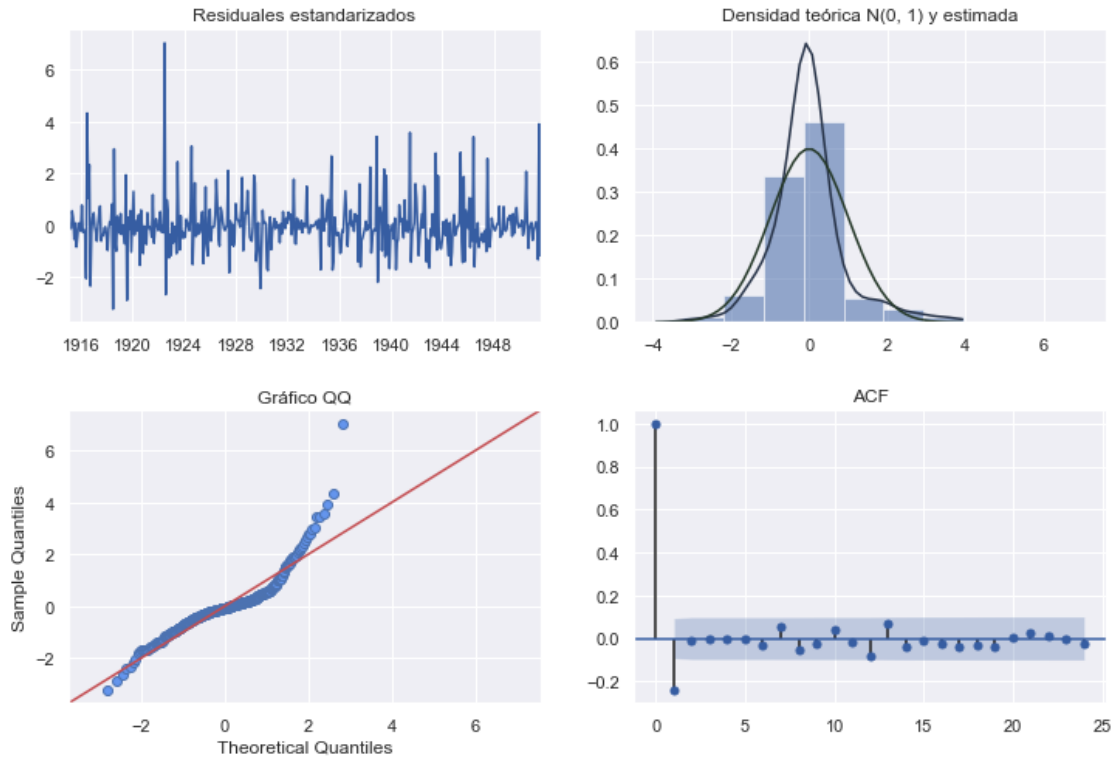
rf_forecast = iter_forecast(rf, y_train, 12, test)

In [140]: rf_undiff_forecast = []
last_obser = train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + rf_forecast[i]
    rf_undiff_forecast.append(yhat)
    last_obser = yhat

rf_undiff_forecast = pd.Series(rf_undiff_forecast)
rf_undiff_forecast.index = test.index

In [141]: resid_diag(rf_resid)
plt.show()

```



```
In [142]: print("Random Forest")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(rf_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(rf_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, rf_undiff_forecast))))
          print("MAE (test): " + str(mae(test, rf_undiff_forecast)))
          print("sMAPE (test): " + str(smape(test, rf_undiff_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 1.2628792392400582e-276

Test de Ljung-Box para k = 6 (p-valor): 0.000340694012737027

Test de Ljung-Box para k = 12 (p-valor): 0.0007856372962224352

```
RMSE (test): 49.99219208276194
MAE (test): 43.52829952097113
sMAPE (test): 74.06849193480416
```

```
In [143]: forecast_plot(test, rf_undiff_forecast)
plt.show()
```



ElasticNet

```
In [144]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:,-1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
        rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

        en_model3 = en.fit(X_val2, y_val2)
```



```

en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], en_forecast))

cvrmse = mean([rmse1, rmse2, rmse3])
if cvrmse < best_score:
    best_score, best_en = cvrmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

```

Best EN: ratio: 0.1 alpha: 0 | CV RMSE = 31.409

```

In [145]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0}

```

```

In [146]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

```

```

In [147]: en_undiff_forecast = []
last_obser = train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + en_forecast[i]
    en_undiff_forecast.append(yhat)
    last_obser = yhat

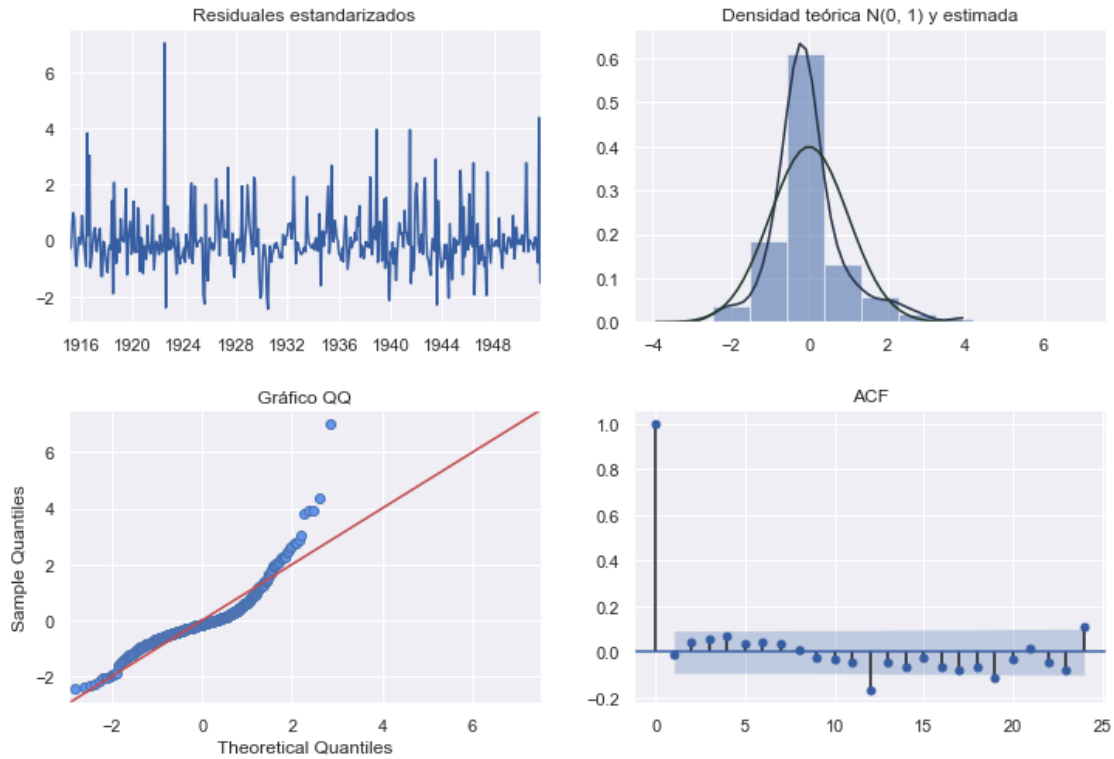
en_undiff_forecast = pd.Series(en_undiff_forecast)
en_undiff_forecast.index = test.index

```

```

In [148]: resid_diag(en_resid)
plt.show()

```



```
In [149]: print("ElasticNet")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(en_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(en_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, en_undiff_forecast))))
          print("MAE (test): " + str(mae(test, en_undiff_forecast)))
          print("sMAPE (test): " + str(smape(test, en_undiff_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Ljung-Box para k = 6 (p-valor): 0.49673055638329544

Test de Ljung-Box para k = 12 (p-valor): 0.07128341852182685

RMSE (test): 28.20851419327389

```
MAE (test): 19.41713375186481
sMAPE (test): 43.647066496865655
```

```
In [150]: forecast_plot(test, en_undiff_forecast)
plt.show()
```



Serie temporal con Box-Cox y desestacionalizada como estructura de DataFrame

```
In [151]: lag_df = series_to_supervised(bc_deseas, n_in=12)
lag_df.head(5)
```

```
Out[151]:
```

	var1(t-12)	var1(t-11)	var1(t-10)	var1(t-9)	var1(t-8)	\
index						
1915-01-01	2.104823	2.032305	2.058984	1.982896	1.945772	
1915-02-01	2.032305	2.058984	1.982896	1.945772	1.870234	
1915-03-01	2.058984	1.982896	1.945772	1.870234	1.894937	
1915-04-01	1.982896	1.945772	1.870234	1.894937	2.008061	
1915-05-01	1.945772	1.870234	1.894937	2.008061	2.104107	

	var1(t-7)	var1(t-6)	var1(t-5)	var1(t-4)	var1(t-3)	var1(t-2)	\
index							
1915-01-01	1.870234	1.894937	2.008061	2.104107	2.094243	2.052516	
1915-02-01	1.894937	2.008061	2.104107	2.094243	2.052516	2.100102	
1915-03-01	2.008061	2.104107	2.094243	2.052516	2.100102	2.040468	
1915-04-01	2.104107	2.094243	2.052516	2.100102	2.040468	1.963905	
1915-05-01	2.094243	2.052516	2.100102	2.040468	1.963905	1.986410	

	var1(t-1)	var1(t)
index		

1915-01-01	2.100102	2.040468
1915-02-01	2.040468	1.963905
1915-03-01	1.963905	1.986410
1915-04-01	1.986410	1.934452
1915-05-01	1.934452	2.072123

Definición de X (predictoras) e y (a predecir)

```
In [152]: X_train = lag_df.iloc[:, :-1]
          y_train = lag_df.iloc[:, -1]
```




Validación cruzada temporal

```
In [153]: train3, validation3 = ts_split(lag_df, train_size=0.75)
          X_train3 = train3.iloc[:, :-1]
          y_train3 = train3.iloc[:, -1]
          X_val3 = validation3.iloc[:, :-1]
          y_val3 = validation3.iloc[:, -1]

          train2, validation2 = ts_split(train3, train_size=0.66)
          X_train2 = train2.iloc[:, :-1]
          y_train2 = train2.iloc[:, -1]
          X_val2 = validation2.iloc[:, :-1]
          y_val2 = validation2.iloc[:, -1]

          train1, validation1 = ts_split(train2, train_size=0.5)
          X_train1 = train1.iloc[:, :-1]
          y_train1 = train1.iloc[:, -1]
          X_val1 = validation1.iloc[:, :-1]
          y_val1 = validation1.iloc[:, -1]

In [154]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
          plt.show()
```



K-Nearest Neighbors

```
In [155]: neighb = list(np.arange(2, 41, 1))
```

```
In [156]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

    print('Best KNN %s CVRMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

```
Best KNN 39 CVRMSE = 0.101
```

```
In [157]: knn_hyp = {"k": 39}
```

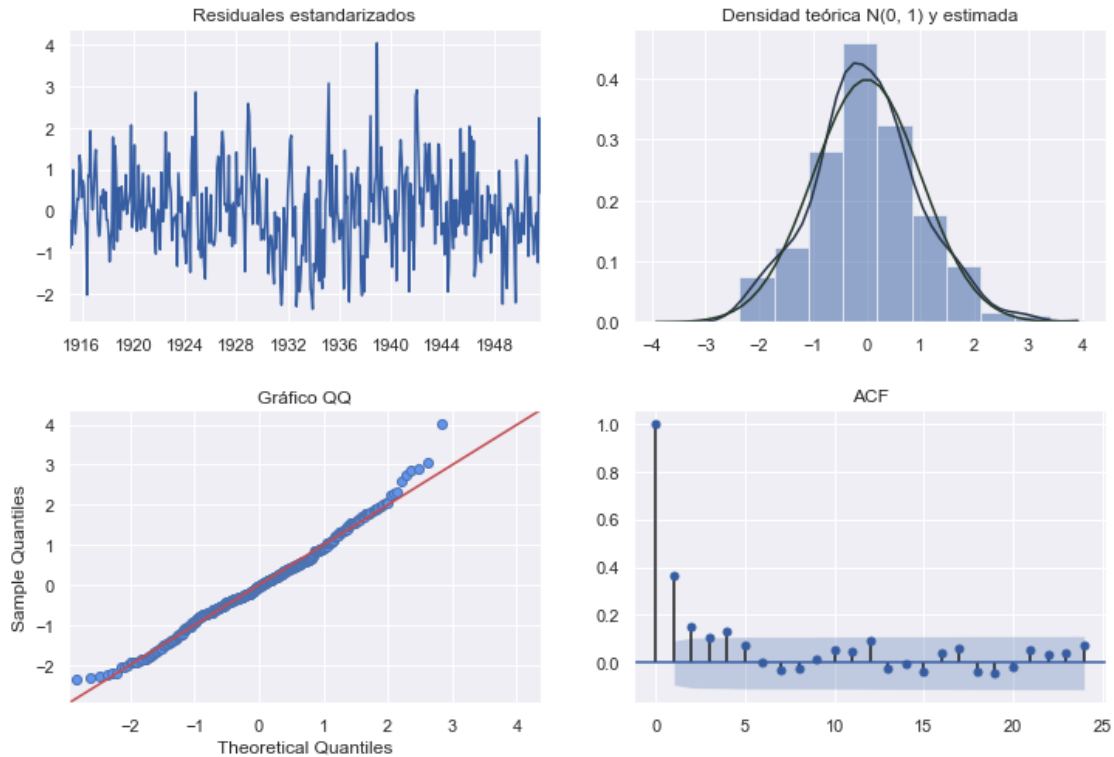
```
In [158]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)
```

```
knn_forecast = iter_forecast(knn, y_train, 12, test)
```

```
In [159]: knn_seas_forecast = knn_forecast + bc_seasonality_forecast
knn_box_forecast = (bc_param * knn_seas_forecast + 1) ** (1 / bc_param)
```

```
In [160]: resid_diag(knn_resid)
plt.show()
```



```
In [161]: print("K-Nearest Neighbors")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(knn_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(knn_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, knn_box_forecast))))
          print("MAE (test): " + str(mae(test, knn_box_forecast)))
          print("sMAPE (test): " + str(smape(test, knn_box_forecast)))
```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 0.0024537980554560092

Test de Ljung-Box para k = 6 (p-valor): 1.7902764790943185e-15

Test de Ljung-Box para k = 12 (p-valor): 7.567142794696897e-14

```
RMSE (test): 33.74267918581206
MAE (test): 20.81188937226639
sMAPE (test): 44.40053551108947
```

```
In [162]: forecast_plot(test, knn_box_forecast)
plt.show()
```



Regression tree

```
In [163]: depths = list(np.arange(1, 21, 1))
```

```
In [164]: best_score, best_tree = float("inf"), None

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)
    tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))

    tree_model2 = tree.fit(X_val1, y_val1)
    tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

    tree_model3 = tree.fit(X_val2, y_val2)
    tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))
```



```

        cvrmse = mean([rmse1, rmse2, rmse3])
        if cvrmse < best_score:
            best_score, best_tree = cvrmse, tree.get_params()

    print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

Best Tree 6 CV RMSE = 0.116

In [165]: tree_hyp = {"depth": 6}

In [166]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                       random_state=seed)
        tree_model = tree.fit(X_train, y_train)
        tree_fitted = pd.Series(tree_model.predict(X_train))
        tree_fitted
        tree_fitted.index = X_train.index
        tree_resid = y_train - tree_fitted
        tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

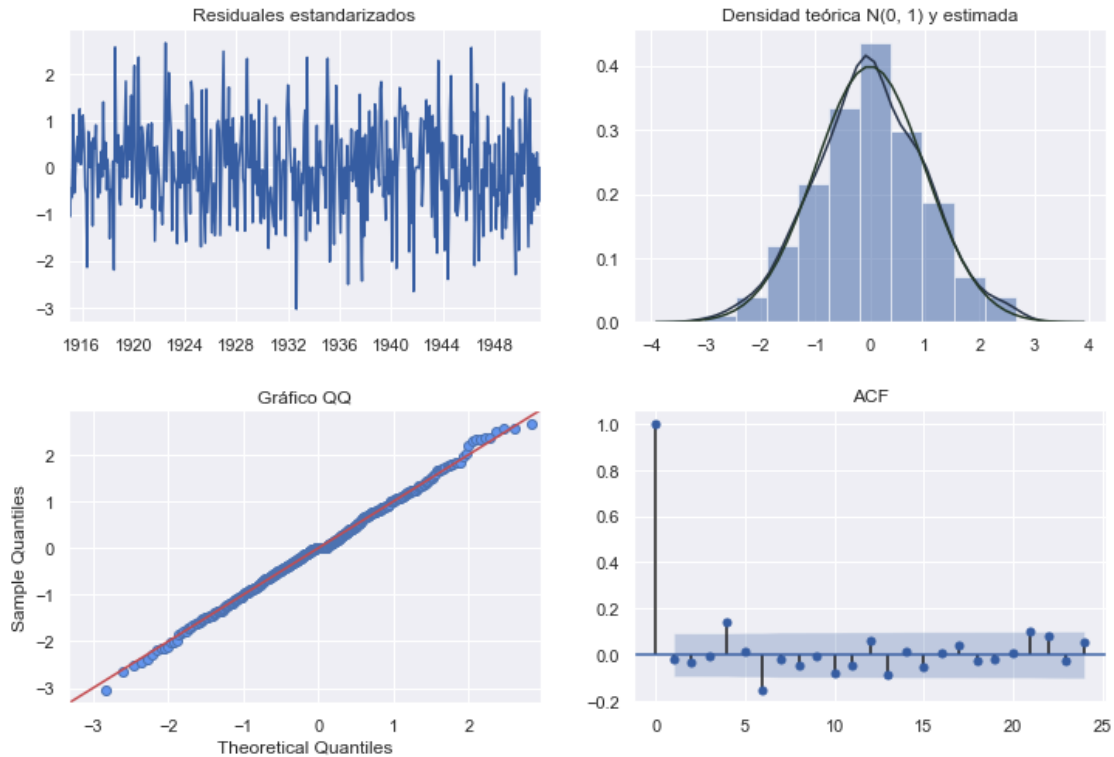
        tree_jb_test = jb(tree_resid)
        tree_lb_test = lb(tree_resid)

        tree_forecast = iter_forecast(tree, y_train, 12, test)

In [167]: tree_seas_forecast = tree_forecast + bc_seasonality_forecast
        tree_box_forecast = (bc_param * tree_seas_forecast + 1) ** (1 / bc_param)

In [168]: resid_diag(tree_resid)
        plt.show()

```



```
In [169]: print("Decision Tree")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(tree_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(tree_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, tree_box_forecast))))
          print("MAE (test): " + str(mae(test, tree_box_forecast)))
          print("sMAPE (test): " + str(smape(test, tree_box_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 0.9929874413510592

Test de Ljung-Box para k = 6 (p-valor): 0.00618239070831623

Test de Ljung-Box para k = 12 (p-valor): 0.005802475974333508

```
RMSE (test): 35.1622944094888
MAE (test): 22.43908436319708
sMAPE (test): 44.0630152274956
```

```
In [170]: forecast_plot(test, tree_box_forecast)
plt.show()
```



Random Forest

```
In [172]: depths = np.linspace(1, 61, 21)
```

```
In [173]: best_score, best_tree = float("inf"), None
```

```
n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]
```

```
for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                      max_depth=depth,
                                      max_features=features,
                                      random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
```

```

rmse1 = sqrt(mse(validation1.iloc[:,-1], rf_forecast))

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f' %
      (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_sc

Best RF: n_estimators: 500 max_depth: 1.0 max_features :1 | CV RMSE = 0.107

In [174]: rf_hyp = {"max_depth": 1,
                  "n_estimators": 500,
                  "max_features": 1}

In [175]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                   n_estimators=rf_hyp["n_estimators"],
                                   max_features=rf_hyp["max_features"],
                                   random_state=seed)

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

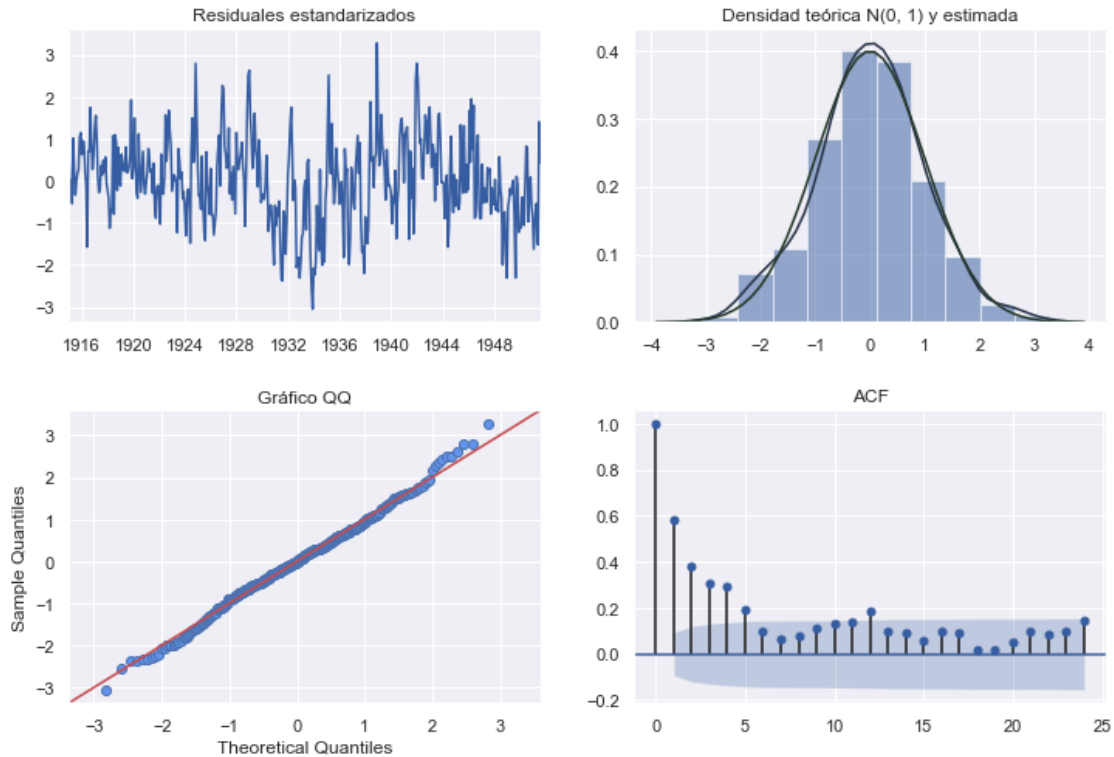
rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

In [176]: rf_seas_forecast = rf_forecast + bc_seasonality_forecast
rf_box_forecast = (bc_param * rf_seas_forecast + 1) ** (1 / bc_param)

In [177]: resid_diag(rf_resid)
plt.show()

```



```
In [178]: print("Random Forest")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(rf_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(rf_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, rf_box_forecast))))
          print("MAE (test): " + str(mae(test, rf_box_forecast)))
          print("sMAPE (test): " + str(smape(test, rf_box_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 0.4770988921732303

Test de Ljung-Box para k = 6 (p-valor): 4.8204864865797315e-65

Test de Ljung-Box para k = 12 (p-valor): 7.749578074539412e-70

```
RMSE (test): 34.36997958512307
MAE (test): 22.24758706621245
sMAPE (test): 44.313099901978816
```

```
In [179]: forecast_plot(test, rf_box_forecast)
plt.show()
```



ElasticNet

```
In [180]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:,-1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
        rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

        en_model3 = en.fit(X_val2, y_val2)
```

```

en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], en_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_en = cvmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

Best EN: ratio: 0.1 alpha: 0 | CV RMSE = 0.106

In [181]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0}

In [182]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                          random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

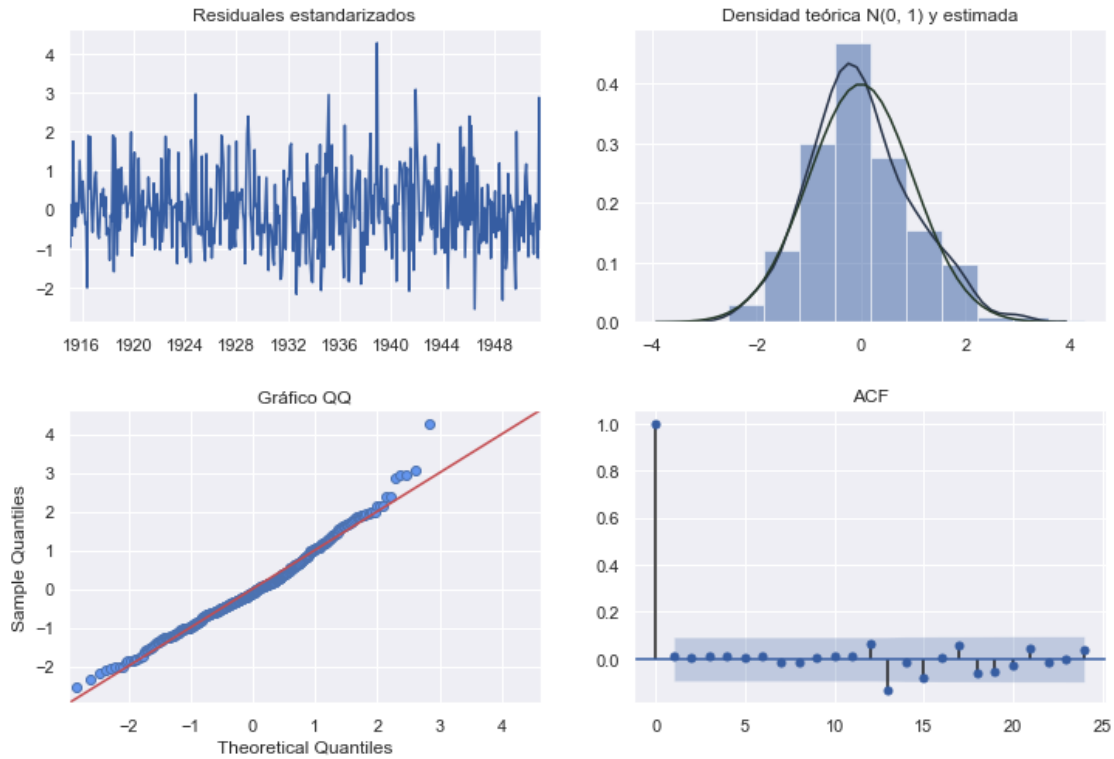
en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

In [183]: en_seas_forecast = en_forecast + bc_seasonality_forecast
en_box_forecast = (bc_param * en_seas_forecast + 1) ** (1 / bc_param)

In [184]: resid_diag(en_resid)
plt.show()

```



```
In [185]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_box_forecast))))
print("MAE (test): " + str(mae(test, en_box_forecast)))
print("sMAPE (test): " + str(smape(test, en_box_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 7.412106994966741e-06

Test de Ljung-Box para k = 6 (p-valor): 0.9996200800570041

Test de Ljung-Box para k = 12 (p-valor): 0.6760069479969495


```


RMSE (test): 34.0536545700252
MAE (test): 21.710813026623878
sMAPE (test): 43.73798854561817

```

```

In [186]: forecast_plot(test, en_box_forecast)
plt.show()

```



Serie temporal con Box-Cox y diferenciada como estructura de DataFrame

```

In [187]: lag_df = series_to_supervised(diff_bc_train, n_in=12)
lag_df.head(5)

```

```

Out[187]:
          var1(t-12)  var1(t-11)  var1(t-10)  var1(t-9)  var1(t-8)  \
index
1915-02-01   -0.022415   -0.036931   -0.120520   -0.050707    0.127983
1915-03-01   -0.036931   -0.120520   -0.050707    0.127983    0.140010
1915-04-01   -0.120520   -0.050707    0.127983    0.140010    0.046348
1915-05-01   -0.050707    0.127983    0.140010    0.046348    0.053055
1915-06-01    0.127983    0.140010    0.046348    0.053055   -0.117011

          var1(t-7)  var1(t-6)  var1(t-5)  var1(t-4)  var1(t-3)  var1(t-2)  \
index
1915-02-01    0.140010    0.046348    0.053055   -0.117011   -0.117013    0.081798
1915-03-01    0.046348    0.053055   -0.117011   -0.117013    0.081798   -0.048950
1915-04-01    0.053055   -0.117011   -0.117013    0.081798   -0.048950   -0.026461
1915-05-01   -0.117011   -0.117013    0.081798   -0.048950   -0.026461   -0.041104
1915-06-01   -0.117013    0.081798   -0.048950   -0.026461   -0.041104   -0.096391

```

	var1(t-1)	var1(t)
index		
1915-02-01	-0.048950	-0.026461
1915-03-01	-0.026461	-0.041104
1915-04-01	-0.041104	-0.096391
1915-05-01	-0.096391	0.124088
1915-06-01	0.124088	0.142337

Definición de X (predictoras) e y (a predecir)

```
In [188]: X_train = lag_df.iloc[:, :-1]
          y_train = lag_df.iloc[:, -1]
```




Validación cruzada temporal

```
In [189]: train3, validation3 = ts_split(lag_df, train_size=0.75)
          X_train3 = train3.iloc[:, :-1]
          y_train3 = train3.iloc[:, -1]
          X_val3 = validation3.iloc[:, :-1]
          y_val3 = validation3.iloc[:, -1]

          train2, validation2 = ts_split(train3, train_size=0.66)
          X_train2 = train2.iloc[:, :-1]
          y_train2 = train2.iloc[:, -1]
          X_val2 = validation2.iloc[:, :-1]
          y_val2 = validation2.iloc[:, -1]

          train1, validation1 = ts_split(train2, train_size=0.5)
          X_train1 = train1.iloc[:, :-1]
          y_train1 = train1.iloc[:, -1]
          X_val1 = validation1.iloc[:, :-1]
          y_val1 = validation1.iloc[:, -1]

In [190]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
          plt.show()
```



K-Nearest Neighbors

```
In [191]: neighb = list(np.arange(2, 41, 1))
```

```
In [192]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

```
Best KNN 21 CVMSE = 0.102
```

```
In [193]: knn_hyp = {"k": 21}
```

```
In [194]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()
```

```
knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)
```

```
knn_forecast = iter_forecast(knn, y_train, 12, test)
```

```
In [195]: knn_undiff_forecast = []
last_obser = bc_train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + knn_forecast[i]
```

```

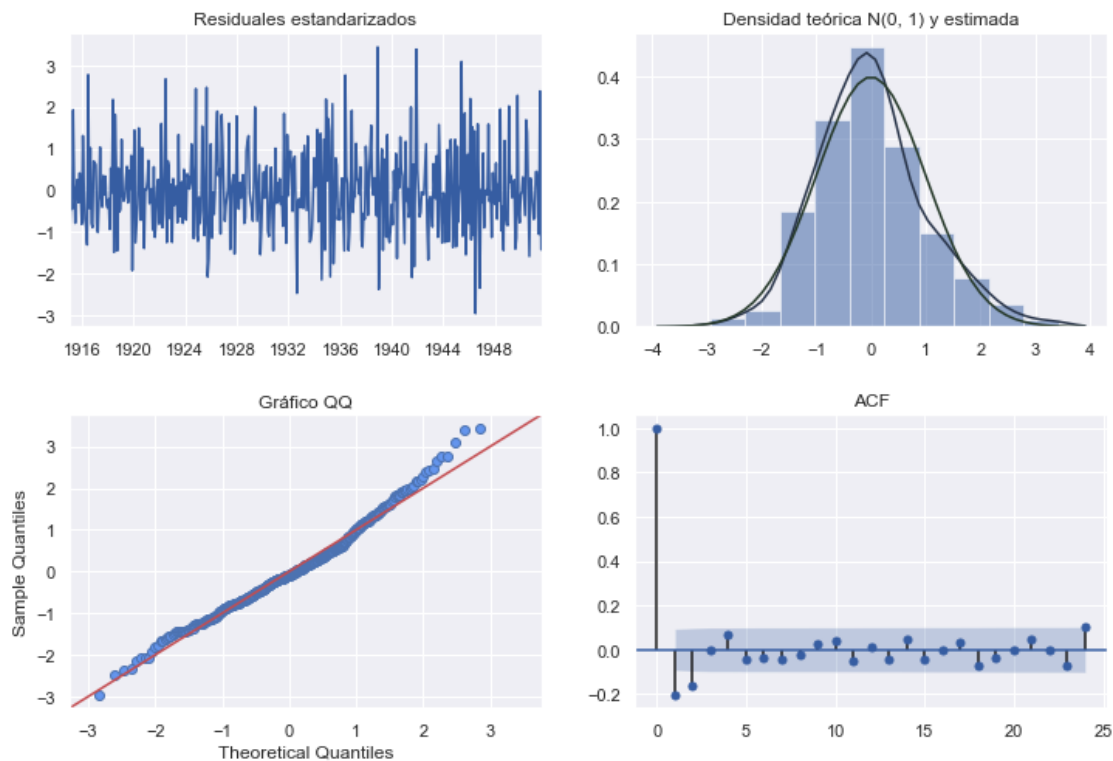
knn_undiff_forecast.append(yhat)
last_obser = yhat

knn_undiff_forecast = pd.Series(knn_undiff_forecast)
knn_undiff_forecast.index = test.index

In [196]: knn_box_forecast = (bc_param * knn_undiff_forecast + 1) ** (1 / bc_param)

In [197]: resid_diag(knn_resid)
plt.show()

```



```

In [198]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " +
      str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")

```

```

print("RMSE (test): " + str(sqrt(mse(test, knn_box_forecast))))
print("MAE (test): " + str(mae(test, knn_box_forecast)))
print("sMAPE (test): " + str(smape(test, knn_box_forecast)))

```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 4.5246797922048964e-05

Test de Ljung-Box para k = 6 (p-valor): 1.426702194406454e-05

Test de Ljung-Box para k = 12 (p-valor): 0.0002938382649062492

RMSE (test): 27.372119226343877


MAE (test): 16.042767653266925

sMAPE (test): 34.04440035728844

```

In [199]: forecast_plot(test, knn_box_forecast)
plt.show()

```



Regression tree

```

In [200]: depths = list(np.arange(1, 21, 1))

```

```

In [201]: best_score, best_tree = float("inf"), None

```

```

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)

```

```

tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:,-1], tree_forecast))

tree_model2 = tree.fit(X_val1, y_val1)
tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], tree_forecast))

tree_model3 = tree.fit(X_val2, y_val2)
tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], tree_forecast))

cvrmse = mean([rmse1, rmse2, rmse3])

if cvrmse < best_score:
    best_score, best_tree = cvrmse, tree.get_params()

print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

```

Best Tree 1 CV RMSE = 0.120

```
In [202]: tree_hyp = {"depth": 1}
```

```
In [203]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                       random_state=seed)
tree_model = tree.fit(X_train, y_train)
tree_fitted = pd.Series(tree_model.predict(X_train))
tree_fitted
tree_fitted.index = X_train.index
tree_resid = y_train - tree_fitted
tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

tree_jb_test = jb(tree_resid)
tree_lb_test = lb(tree_resid)

```

```
tree_forecast = iter_forecast(tree, y_train, 12, test)
```

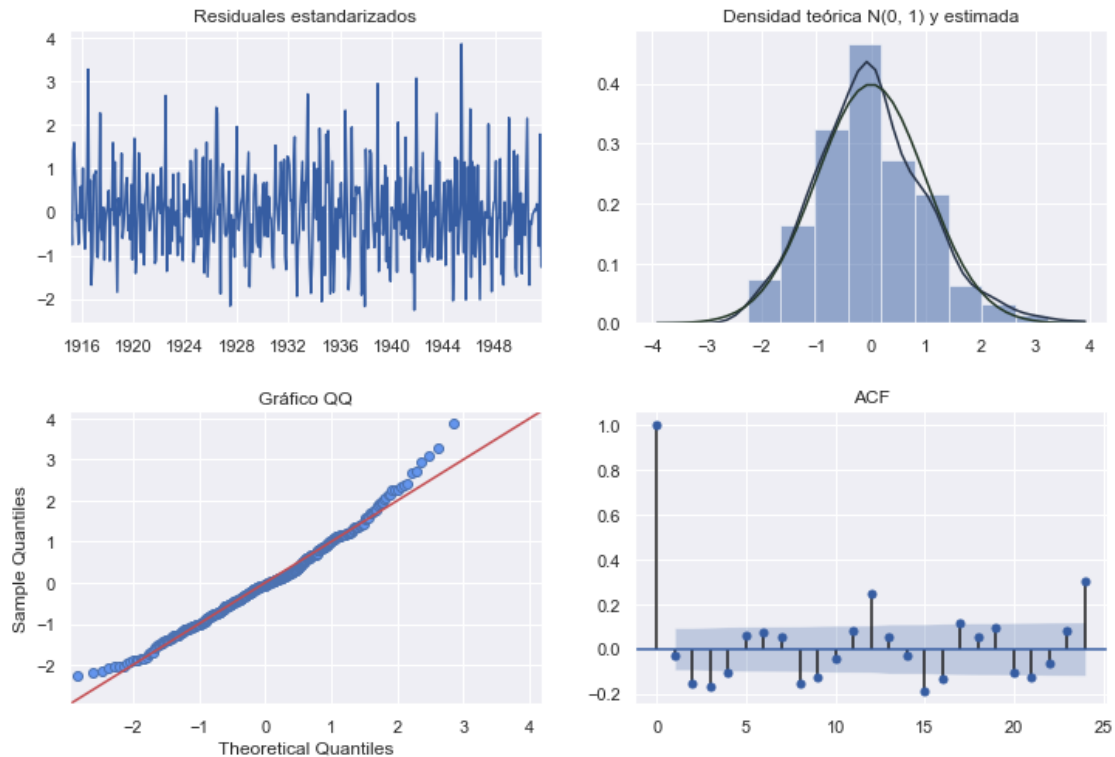
```
In [204]: tree_undiff_forecast = []
last_obser = bc_train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + tree_forecast[i]
    tree_undiff_forecast.append(yhat)
    last_obser = yhat

tree_undiff_forecast = pd.Series(tree_undiff_forecast)
tree_undiff_forecast.index = test.index

```

```
In [205]: tree_box_forecast = (bc_param * tree_undiff_forecast + 1) ** (1 / bc_param)
```

```
In [206]: resid_diag(tree_resid)
plt.show()
```



```
In [207]: print("Decision Tree")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " +
      str(tree_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(tree_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(tree_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, tree_box_forecast))))
print("MAE (test): " + str(mae(test, tree_box_forecast)))
print("sMAPE (test): " + str(smape(test, tree_box_forecast)))
```


Decision Tree

Observaciones ajustadas: 452
Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 6.676127468555702e-05
Test de Ljung-Box para k = 6 (p-valor): 2.6885859567584192e-05
Test de Ljung-Box para k = 12 (p-valor): 1.6607113930256373e-12

RMSE (test): 30.862439905297794
MAE (test): 19.343326550124953
sMAPE (test): 41.765027713266996

```
In [208]: forecast_plot(test, tree_box_forecast)
plt.show()
```



Random Forest

```
In [209]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                       max_depth=depth,
                                       max_features=features,
                                       random_state=seed)
```

```

rf_model1 = rf.fit(X_train1, y_train1)
rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f'
      % (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_

Best RF: n_estimators: 500 max_depth: 4.157894736842105 max_features :0.8 | CV RMSE = 0.113

In [210]: rf_hyp = {"max_depth": 4.15,
                    "n_estimators": 500,
                    "max_features": 0.8}

In [211]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                     n_estimators=rf_hyp["n_estimators"],
                                     max_features=rf_hyp["max_features"],
                                     random_state=seed)

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

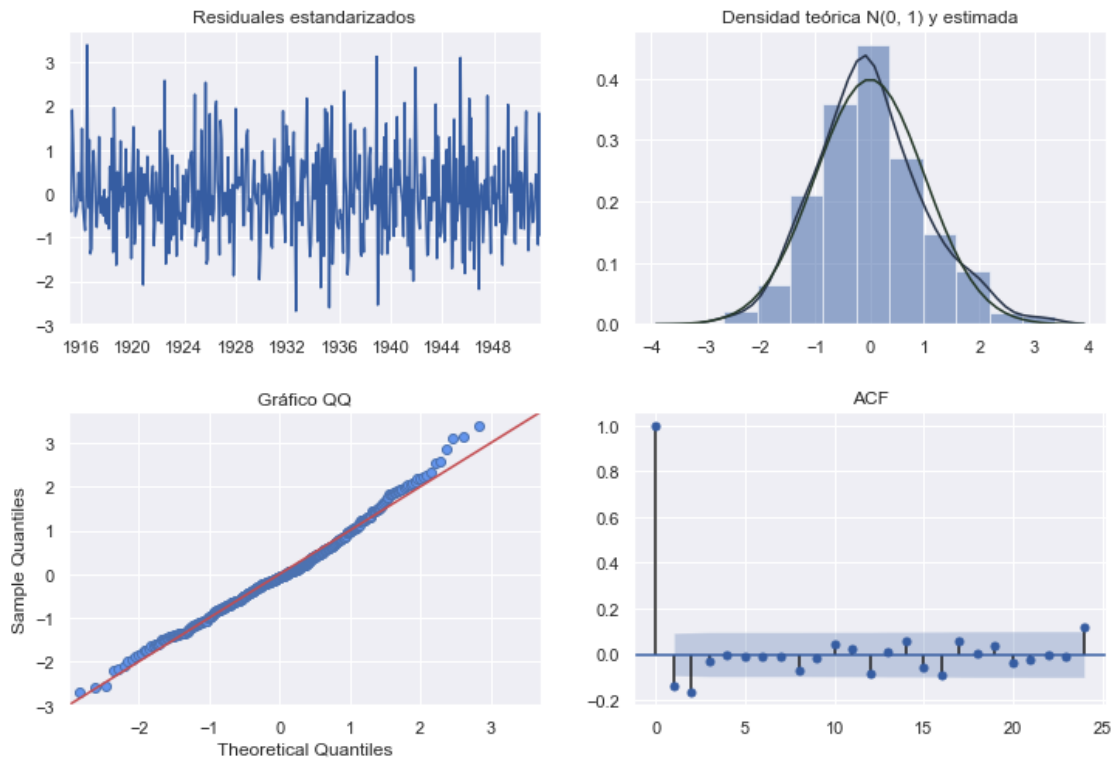
In [212]: rf_undiff_forecast = []
last_obser = bc_deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + rf_forecast[i]
    rf_undiff_forecast.append(yhat)
    last_obser = yhat

```

```
rf_undiff_forecast = pd.Series(rf_undiff_forecast)
rf_undiff_forecast.index = test.index
```

```
In [213]: rf_box_forecast = (bc_param * rf_undiff_forecast + 1) ** (1 / bc_param)
```

```
In [214]: resid_diag(rf_resid)
plt.show()
```



```
In [215]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_box_forecast))))
print("MAE (test): " + str(mae(test, rf_box_forecast)))
print("sMAPE (test): " + str(smape(test, rf_box_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 0.0017866076588243709

Test de Ljung-Box para k = 6 (p-valor): 0.0036130029610332488


Test de Ljung-Box para k = 12 (p-valor): 0.008784333768246519

RMSE (test): 43.79388980909914

MAE (test): 29.732433883121928

sMAPE (test): 90.94202211425733

```
In [216]: forecast_plot(test, rf_box_forecast)
plt.show()
```



ElasticNet

```
In [217]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
```

```

en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:,-1], en_forecast))

en_model2 = en.fit(X_val1, y_val1)
en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

en_model3 = en.fit(X_val2, y_val2)
en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], en_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_en = cvmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

```

Best EN: ratio: 0.1 alpha: 0 | CV RMSE = 0.111

```

In [218]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0}

```

```

In [219]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

```

```

In [220]: en_undiff_forecast = []
last_obser = bc_train.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + en_forecast[i]
    en_undiff_forecast.append(yhat)
    last_obser = yhat

en_undiff_forecast = pd.Series(en_undiff_forecast)
en_undiff_forecast.index = test.index

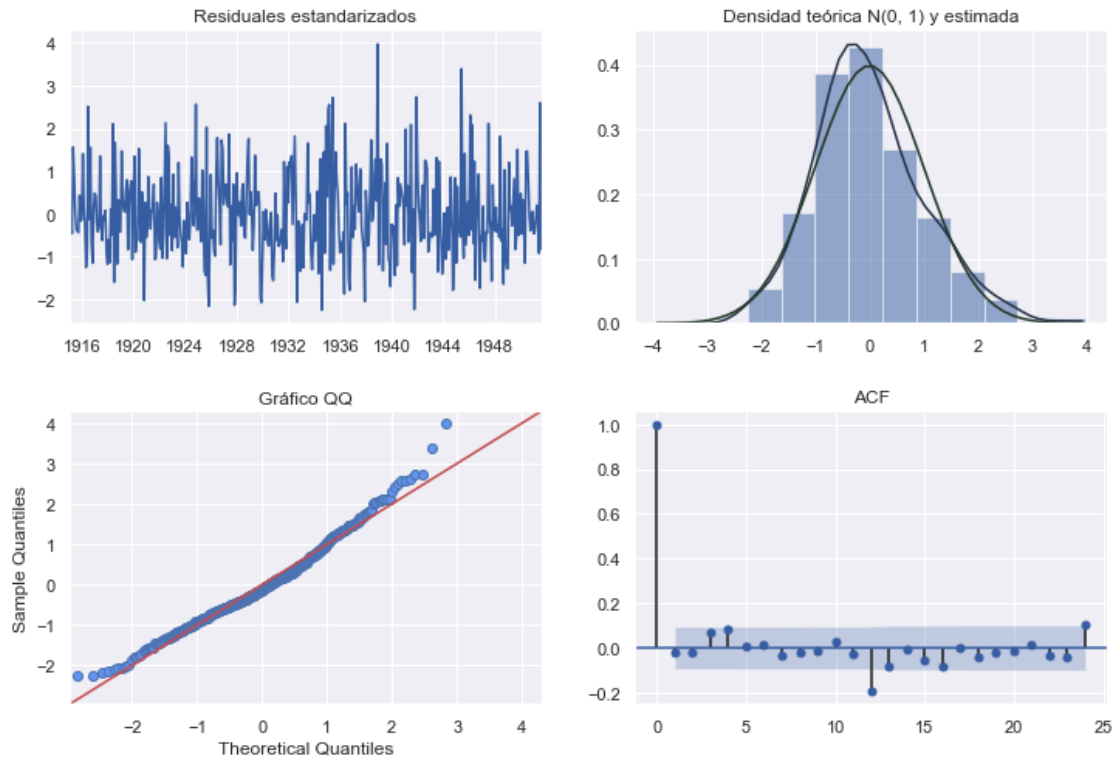
```

```

In [221]: en_box_forecast = (bc_param * en_undiff_forecast + 1) ** (1 / bc_param)

```

```
In [222]: resid_diag(en_resid)
plt.show()
```



```
In [223]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_box_forecast))))
print("MAE (test): " + str(mae(test, en_box_forecast)))
print("sMAPE (test): " + str(smape(test, en_box_forecast)))
```


ElasticNet

```
Observaciones ajustadas: 452
Observaciones predichas: 112
```

Test de Jarque-Bera (p-valor): 1.621113733273755e-06
 Test de Ljung-Box para k = 6 (p-valor): 0.5239960585921445
 Test de Ljung-Box para k = 12 (p-valor): 0.012392484507714213

RMSE (test): 33.327165564718044
 MAE (test): 18.778881608977922
 sMAPE (test): 39.585127561297675

```
In [224]: forecast_plot(test, en_box_forecast)
plt.show()
```



Serie temporal desestacionalizada y diferenciada como estructura de DataFrame

```
In [225]: diff_deseas = deseas.diff(1).iloc[1:]
```

```
In [226]: lag_df = series_to_supervised(diff_deseas, n_in=12)
lag_df.head(5)
```

```
Out[226]:
```

	var1(t-12)	var1(t-11)	var1(t-10)	var1(t-9)	var1(t-8)	\
index						
1915-02-01	-12.568751	4.257272	-10.102387	-2.545743	-27.175797	
1915-03-01	4.257272	-10.102387	-2.545743	-27.175797	-17.669263	
1915-04-01	-10.102387	-2.545743	-27.175797	-17.669263	54.286194	
1915-05-01	-2.545743	-27.175797	-17.669263	54.286194	39.474457	
1915-06-01	-27.175797	-17.669263	54.286194	39.474457	-25.286546	

	var1(t-7)	var1(t-6)	var1(t-5)	var1(t-4)	var1(t-3)	var1(t-2)	\
index							
1915-02-01	-17.669263	54.286194	39.474457	-25.286546	-14.752516	8.114108	
1915-03-01	54.286194	39.474457	-25.286546	-14.752516	8.114108	-9.567028	

1915-04-01	39.474457	-25.286546	-14.752516	8.114108	-9.567028	-11.491751
1915-05-01	-25.286546	-14.752516	8.114108	-9.567028	-11.491751	6.097272
1915-06-01	-14.752516	8.114108	-9.567028	-11.491751	6.097272	-3.561387

	var1(t-1)	var1(t)
index		
1915-02-01	-9.567028	-11.491751
1915-03-01	-11.491751	6.097272
1915-04-01	6.097272	-3.561387
1915-05-01	-3.561387	14.444257
1915-06-01	14.444257	-9.590797

Definición de X (predictoras) e y (a predecir)

```
In [227]: X_train = lag_df.iloc[:, :-1]
          y_train = lag_df.iloc[:, -1]
```




Validación cruzada temporal

```
In [228]: train3, validation3 = ts_split(lag_df, train_size=0.75)
          X_train3 = train3.iloc[:, :-1]
          y_train3 = train3.iloc[:, -1]
          X_val3 = validation3.iloc[:, :-1]
          y_val3 = validation3.iloc[:, -1]

          train2, validation2 = ts_split(train3, train_size=0.66)
          X_train2 = train2.iloc[:, :-1]
          y_train2 = train2.iloc[:, -1]
          X_val2 = validation2.iloc[:, :-1]
          y_val2 = validation2.iloc[:, -1]

          train1, validation1 = ts_split(train2, train_size=0.5)
          X_train1 = train1.iloc[:, :-1]
          y_train1 = train1.iloc[:, -1]
          X_val1 = validation1.iloc[:, :-1]
          y_val1 = validation1.iloc[:, -1]

In [229]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
          plt.show()
```

K-Nearest Neighbors

```
In [230]: neighb = list(np.arange(2, 41, 1))
```

```
In [231]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

Best KNN 33 CVMSE = 25.992

```
In [232]: knn_hyp = {"k": 33}
```

```
In [233]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)
```

```
knn_forecast = iter_forecast(knn, y_train, 12, test)
```

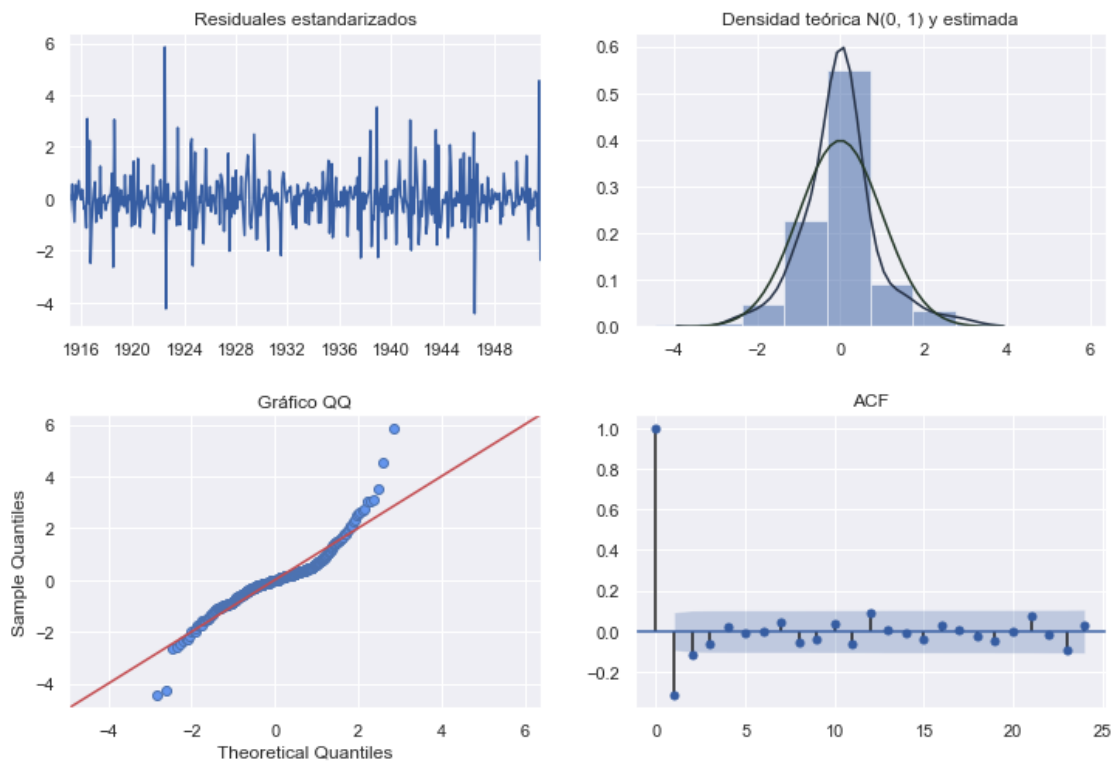
```
In [234]: knn_undiff_forecast = []
last_obser = deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + knn_forecast[i]
```

```
knn_undiff_forecast.append(yhat)
last_obser = yhat
```

```
knn_undiff_forecast = pd.Series(knn_undiff_forecast)
knn_undiff_forecast.index = test.index
```

```
In [235]: knn_seas_forecast = knn_undiff_forecast + seasonality_forecast
```

```
In [236]: resid_diag(knn_resid)
plt.show()
```



```
In [237]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, knn_seas_forecast))))
```

```
print("MAE (test): " + str(mae(test, knn_seas_forecast)))
print("sMAPE (test): " + str(smape(test, knn_seas_forecast)))
```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 5.717097260166605e-124

Test de Ljung-Box para k = 6 (p-valor): 7.070382754371443e-09


Test de Ljung-Box para k = 12 (p-valor): 5.902152283903449e-08

RMSE (test): 84.05832496249754

MAE (test): 63.474552623087746

sMAPE (test): 132.40169832978398

```
In [238]: forecast_plot(test, knn_seas_forecast)
plt.show()
```



Regresion tree

```
In [239]: depths = list(np.arange(1, 21, 1))
```

```
In [240]: best_score, best_tree = float("inf"), None
```

```
for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)
    tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
```

```

rmse1 = sqrt(mse(validation1.iloc[:, -1], tree_forecast))

tree_model2 = tree.fit(X_val1, y_val1)
tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], tree_forecast))

tree_model3 = tree.fit(X_val2, y_val2)
tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], tree_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_tree = cvmse, tree.get_params()

print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

```

Best Tree 1 CV RMSE = 26.466

```
In [241]: tree_hyp = {"depth": 1}
```

```
In [242]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                     random_state=seed)

tree_model = tree.fit(X_train, y_train)
tree_fitted = pd.Series(tree_model.predict(X_train))
tree_fitted
tree_fitted.index = X_train.index
tree_resid = y_train - tree_fitted
tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

tree_jb_test = jb(tree_resid)
tree_lb_test = lb(tree_resid)

tree_forecast = iter_forecast(tree, y_train, 12, test)

```

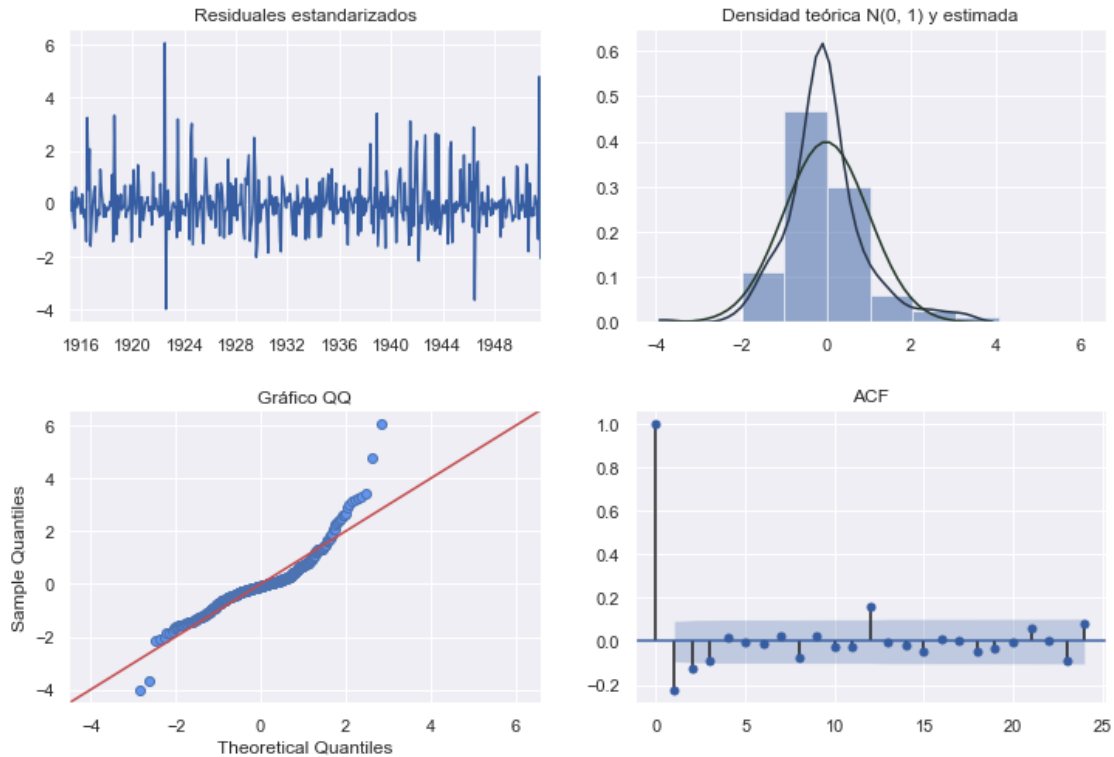
```
In [243]: tree_undiff_forecast = []
last_obser = deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + tree_forecast[i]
    tree_undiff_forecast.append(yhat)
    last_obser = yhat

tree_undiff_forecast = pd.Series(tree_undiff_forecast)
tree_undiff_forecast.index = test.index

```

```
In [244]: tree_seas_forecast = tree_undiff_forecast + seasonality_forecast
```

```
In [245]: resid_diag(tree_resid)
plt.show()
```



```
In [246]: print("Decision Tree")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(tree_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(tree_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, tree_seas_forecast))))
          print("MAE (test): " + str(mae(test, tree_seas_forecast)))
          print("sMAPE (test): " + str(smape(test, tree_seas_forecast)))
```

Decision Tree

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 5.984587735691505e-150

Test de Ljung-Box para k = 6 (p-valor): 3.310821158172319e-05

Test de Ljung-Box para k = 12 (p-valor): 7.717590622639042e-06

```
RMSE (test): 220.1524782212875
MAE (test): 192.52545166614178
sMAPE (test): 130.60223449593917
```

```
In [247]: forecast_plot(test, tree_seas_forecast)
plt.show()
```



Random Forest

```
In [248]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                      max_depth=depth,
                                      max_features=features,
                                      random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
            rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
            rmse1 = sqrt(mse(validation1.iloc[:, -1], rf_forecast))
```

```

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:, -1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:, -1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f' %
      (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_score))

Best RF: n_estimators: 500 max_depth: 1.0 max_features :1 | CV RMSE = 25.879

In [249]: rf_hyp = {"max_depth": 1,
                  "n_estimators": 500,
                  "max_features": 1}

In [250]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                     n_estimators=rf_hyp["n_estimators"],
                                     max_features=rf_hyp["max_features"],
                                     random_state=seed)

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

In [251]: rf_undiff_forecast = []
last_obser = deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + rf_forecast[i]
    rf_undiff_forecast.append(yhat)
    last_obser = yhat

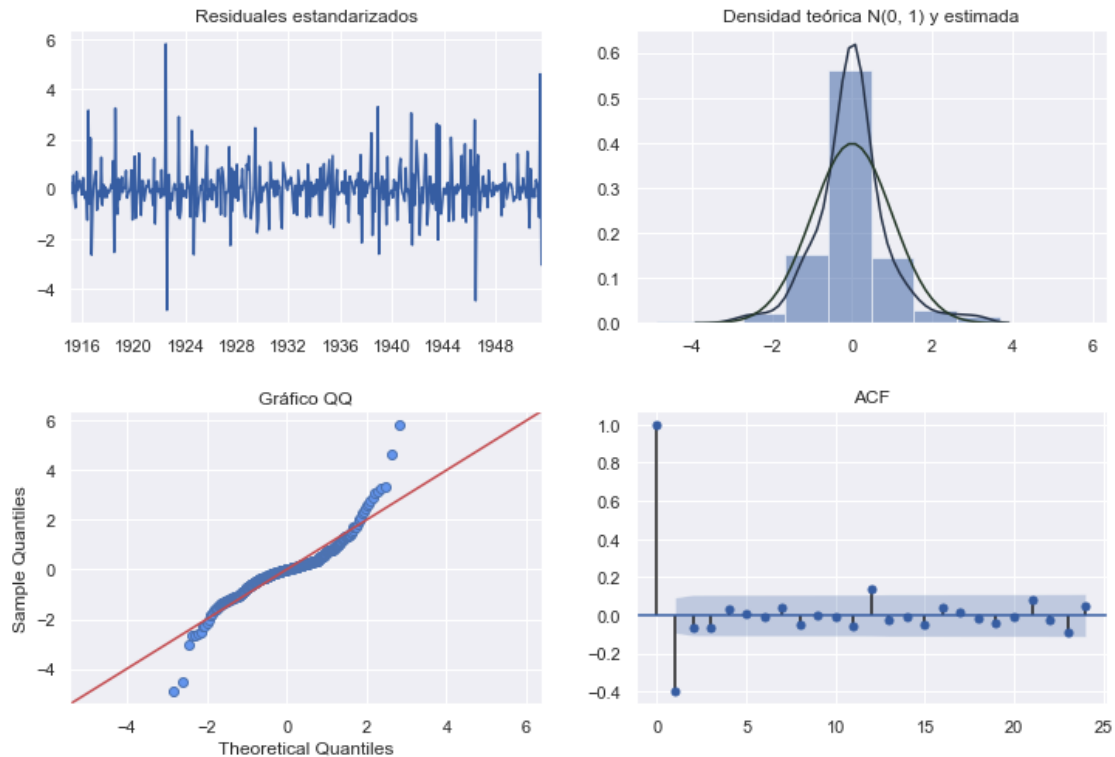
rf_undiff_forecast = pd.Series(rf_undiff_forecast)
rf_undiff_forecast.index = test.index

In [252]: rf_seas_forecast = rf_undiff_forecast + seasonality_forecast

```



```
In [253]: resid_diag(rf_resid)
plt.show()
```



```
In [254]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_seas_forecast))))
print("MAE (test): " + str(mae(test, rf_seas_forecast)))
print("sMAPE (test): " + str(smape(test, rf_seas_forecast)))
```


Random Forest

Observaciones ajustadas: 452
Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 1.236348856640542e-154
 Test de Ljung-Box para k = 6 (p-valor): 2.0612833395810488e-13
 Test de Ljung-Box para k = 12 (p-valor): 9.123411363294815e-13

RMSE (test): 39.87890461677487
 MAE (test): 29.491078404147572
 sMAPE (test): 56.01068915026782

```
In [255]: forecast_plot(test, rf_seas_forecast)
plt.show()
```



ElasticNet

```
In [256]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
        en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
        rmse1 = sqrt(mse(validation1.iloc[:, -1], en_forecast))

        en_model2 = en.fit(X_val1, y_val1)
        en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
```

```

rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

en_model3 = en.fit(X_val2, y_val2)
en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], en_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_en = cvmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

```

Best EN: ratio: 0.1 alpha: 0 | CV RMSE = 25.778

```

In [257]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0}

```

```

In [258]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

```

```

In [259]: en_undiff_forecast = []
last_obser = deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + en_forecast[i]
    en_undiff_forecast.append(yhat)
    last_obser = yhat

```

```

en_undiff_forecast = pd.Series(en_undiff_forecast)
en_undiff_forecast.index = test.index

```

```

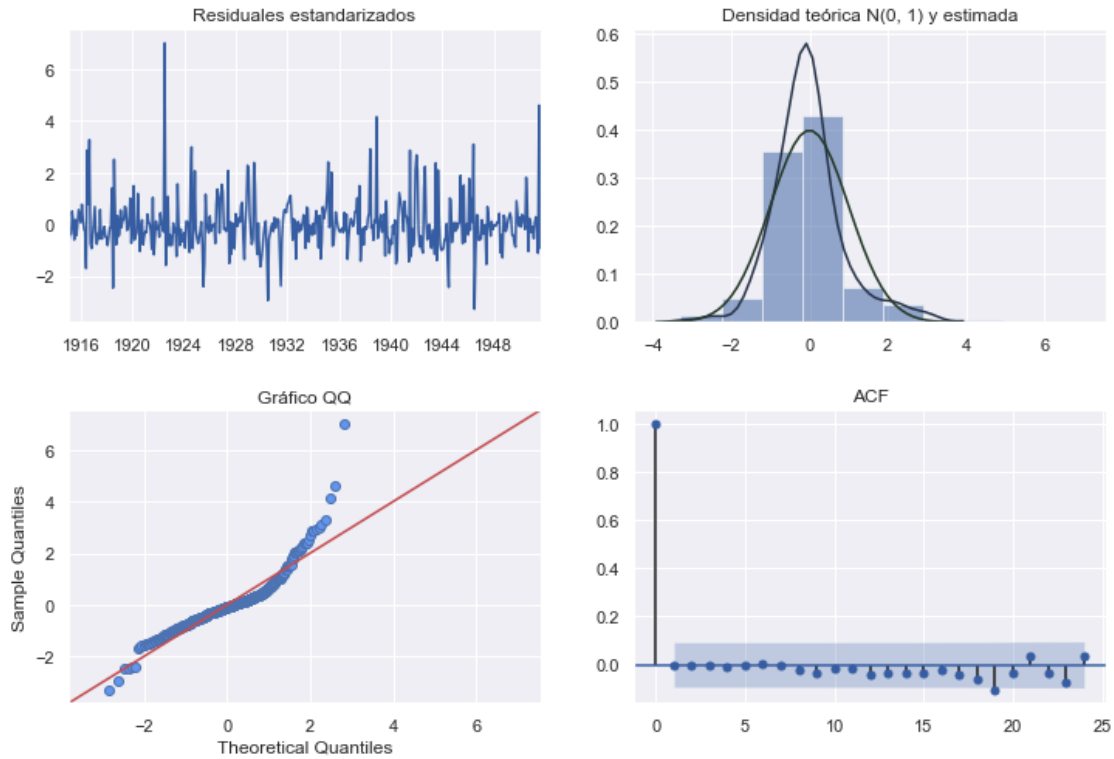
In [260]: en_seas_forecast = en_undiff_forecast + seasonality_forecast

```

```

In [261]: resid_diag(en_resid)
plt.show()

```



```
In [262]: print("ElasticNet")
          print("")
          print("Observaciones ajustadas: " + str(len(train)))
          print("Observaciones predichas: " + str(len(test)))
          print("")
          print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
          print("Test de Ljung-Box para k = 6 (p-valor): " +
                str(en_lb_test[1][6]))
          print("Test de Ljung-Box para k = 12 (p-valor): " +
                str(en_lb_test[1][12]))
          print("")
          print("RMSE (test): " + str(sqrt(mse(test, en_seas_forecast))))
          print("MAE (test): " + str(mae(test, en_seas_forecast)))
          print("sMAPE (test): " + str(smape(test, en_seas_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112


Test de Jarque-Bera (p-valor): 2.0755161120405965e-264

Test de Ljung-Box para k = 6 (p-valor): 0.9999999468181409

Test de Ljung-Box para k = 12 (p-valor): 0.9994838286140066

```
RMSE (test): 35.74647450822158
MAE (test): 24.364237928781076
sMAPE (test): 47.541746791674875
```

```
In [263]: forecast_plot(test, en_seas_forecast)
plt.show()
```



Serie temporal con Box-Cox, desestacionalizada y diferenciada como estructura de DataFrame

```
In [264]: lag_df = series_to_supervised(bc_diff_deseas, n_in=12)
lag_df.head(5)
```

```
Out[264]:
```

	var1(t-12)	var1(t-11)	var1(t-10)	var1(t-9)	var1(t-8)	\
index						
1915-02-01	-0.072518	0.026679	-0.076088	-0.037124	-0.075538	
1915-03-01	0.026679	-0.076088	-0.037124	-0.075538	0.024702	
1915-04-01	-0.076088	-0.037124	-0.075538	0.024702	0.113124	
1915-05-01	-0.037124	-0.075538	0.024702	0.113124	0.096046	
1915-06-01	-0.075538	0.024702	0.113124	0.096046	-0.009864	

	var1(t-7)	var1(t-6)	var1(t-5)	var1(t-4)	var1(t-3)	var1(t-2)	\
index							
1915-02-01	0.024702	0.113124	0.096046	-0.009864	-0.041726	0.047586	
1915-03-01	0.113124	0.096046	-0.009864	-0.041726	0.047586	-0.059635	
1915-04-01	0.096046	-0.009864	-0.041726	0.047586	-0.059635	-0.076563	
1915-05-01	-0.009864	-0.041726	0.047586	-0.059635	-0.076563	0.022506	
1915-06-01	-0.041726	0.047586	-0.059635	-0.076563	0.022506	-0.051959	

	var1(t-1)	var1(t)
index		
1915-02-01	-0.059635	-0.076563
1915-03-01	-0.076563	0.022506
1915-04-01	0.022506	-0.051959
1915-05-01	-0.051959	0.137671
1915-06-01	0.137671	-0.061183

Definición de X (predictoras) e y (a predecir)

```
In [265]: X_train = lag_df.iloc[:, :-1]
          y_train = lag_df.iloc[:, -1]
```




Validación cruzada temporal

```
In [266]: train3, validation3 = ts_split(lag_df, train_size=0.75)
          X_train3 = train3.iloc[:, :-1]
          y_train3 = train3.iloc[:, -1]
          X_val3 = validation3.iloc[:, :-1]
          y_val3 = validation3.iloc[:, -1]

          train2, validation2 = ts_split(train3, train_size=0.66)
          X_train2 = train2.iloc[:, :-1]
          y_train2 = train2.iloc[:, -1]
          X_val2 = validation2.iloc[:, :-1]
          y_val2 = validation2.iloc[:, -1]

          train1, validation1 = ts_split(train2, train_size=0.5)
          X_train1 = train1.iloc[:, :-1]
          y_train1 = train1.iloc[:, -1]
          X_val1 = validation1.iloc[:, :-1]
          y_val1 = validation1.iloc[:, -1]

In [267]: tscv_plot(train1["var1(t)"], validation1["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train2["var1(t)"], validation2["var1(t)"], train.index[-1])
          plt.show()
          tscv_plot(train3["var1(t)"], validation3["var1(t)"], train.index[-1])
          plt.show()
```



K-Nearest Neighbors

```
In [268]: neighb = list(np.arange(2, 41, 1))
```

```
In [269]: best_score, best_knn = float("inf"), None
```

```
for k in neighb:
    knn = KNeighborsRegressor(n_neighbors=k)

    knn_model1 = knn.fit(X_train1, y_train1)
    knn_forecast = iter_forecast(knn_model1, y_train1, 12, validation1)
    rmse1 = sqrt(mse(validation1.iloc[:, -1], knn_forecast))

    knn_model2 = knn.fit(X_val1, y_val1)
    knn_forecast = iter_forecast(knn_model2, y_val1, 12, validation2)
    rmse2 = sqrt(mse(validation2.iloc[:, -1], knn_forecast))

    knn_model3 = knn.fit(X_val2, y_val2)
    knn_forecast = iter_forecast(knn_model3, y_val2, 12, validation3)
    rmse3 = sqrt(mse(validation3.iloc[:, -1], knn_forecast))

    cvrmse = mean([rmse1, rmse2, rmse3])
    if cvrmse < best_score:
        best_score, best_knn = cvrmse, knn.get_params()

print('Best KNN %s CVRMSE = %.3f' % (best_knn['n_neighbors'], best_score))
```

```
Best KNN 27 CVRMSE = 0.084
```

```
In [270]: knn_hyp = {"k": 27}
```

```
In [271]: knn = KNeighborsRegressor(n_neighbors=knn_hyp["k"])
knn_model = knn.fit(X_train, y_train)
knn_fitted = pd.Series(knn_model.predict(X_train))
knn_fitted
knn_fitted.index = X_train.index
knn_resid = y_train - knn_fitted
knn_resid = (knn_resid - knn_resid.mean()) / knn_resid.std()

knn_jb_test = jb(knn_resid)
knn_lb_test = lb(knn_resid)
```

```
knn_forecast = iter_forecast(knn, y_train, 12, test)
```

```
In [272]: knn_undiff_forecast = []
last_obser = bc_deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + knn_forecast[i]
```



```

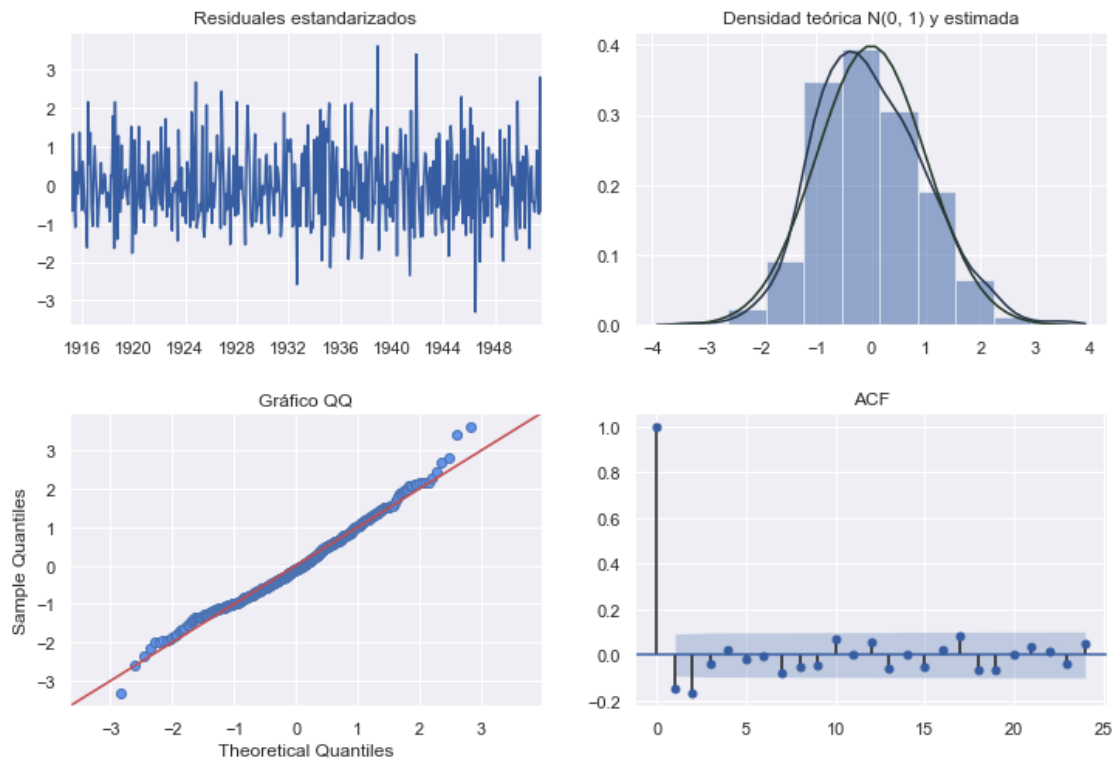
knn_undiff_forecast.append(yhat)
last_obser = yhat

knn_undiff_forecast = pd.Series(knn_undiff_forecast)
knn_undiff_forecast.index = test.index

In [273]: knn_seas_forecast = knn_undiff_forecast + bc_seasonality_forecast
knn_box_forecast = (bc_param * knn_seas_forecast + 1) ** (1 / bc_param)

In [274]: resid_diag(knn_resid)
plt.show()

```



```

In [275]: print("K-Nearest Neighbors")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(knn_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(knn_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(knn_lb_test[1][12]))
print("")

```

```

print("RMSE (test): " + str(sqrt(mse(test, knn_box_forecast))))
print("MAE (test): " + str(mae(test, knn_box_forecast)))
print("sMAPE (test): " + str(smape(test, knn_box_forecast)))

```

K-Nearest Neighbors

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 0.003383083593474169

Test de Ljung-Box para k = 6 (p-valor): 0.0005375519245986289

Test de Ljung-Box para k = 12 (p-valor): 0.0015617704092770694

RMSE (test): 36.49420211710823


MAE (test): 22.4493403519271

sMAPE (test): 52.483487500997285

```

In [276]: forecast_plot(test, knn_box_forecast)
plt.show()

```



Regression tree

```

In [277]: depths = list(np.arange(1, 21, 1))

```

```

In [278]: best_score, best_tree = float("inf"), None

```

```

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=seed)

    tree_model1 = tree.fit(X_train1, y_train1)

```

```

tree_forecast = iter_forecast(tree_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:,-1], tree_forecast))

tree_model2 = tree.fit(X_val1, y_val1)
tree_forecast = iter_forecast(tree_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], tree_forecast))

tree_model3 = tree.fit(X_val2, y_val2)
tree_forecast = iter_forecast(tree_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], tree_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_tree = cvmse, tree.get_params()

print('Best Tree %s CV RMSE = %.3f' % (best_tree['max_depth'], best_score))

```

Best Tree 1 CV RMSE = 0.088

```
In [279]: tree_hyp = {"depth": 1}
```

```
In [280]: tree = DecisionTreeRegressor(max_depth=tree_hyp["depth"],
                                     random_state=seed)

tree_model = tree.fit(X_train, y_train)
tree_fitted = pd.Series(tree_model.predict(X_train))
tree_fitted
tree_fitted.index = X_train.index
tree_resid = y_train - tree_fitted
tree_resid = (tree_resid - tree_resid.mean()) / tree_resid.std()

tree_jb_test = jb(tree_resid)
tree_lb_test = lb(tree_resid)

tree_forecast = iter_forecast(tree, y_train, 12, test)

```

```
In [281]: tree_undiff_forecast = []
last_obser = bc_deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + tree_forecast[i]
    tree_undiff_forecast.append(yhat)
    last_obser = yhat

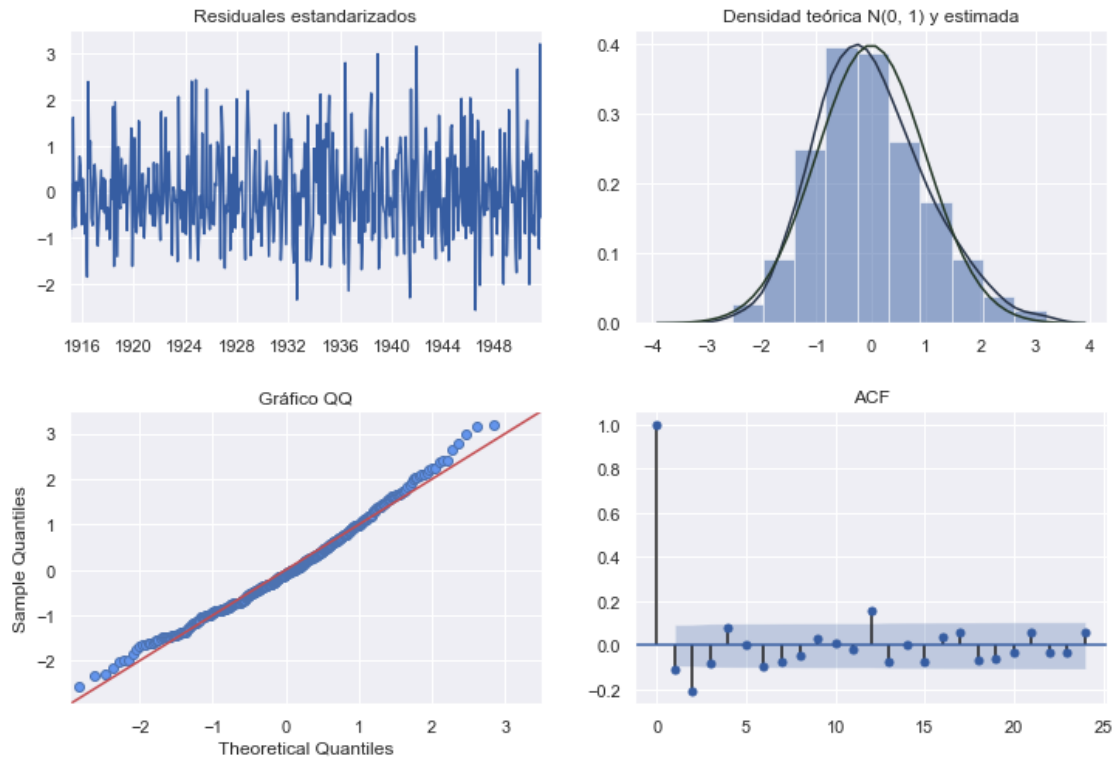
tree_undiff_forecast = pd.Series(tree_undiff_forecast)
tree_undiff_forecast.index = test.index

```

```
In [282]: tree_seas_forecast = knn_undiff_forecast + bc_seasonality_forecast
tree_box_forecast = (bc_param * tree_seas_forecast + 1) ** (1 / bc_param)

```

```
In [283]: resid_diag(tree_resid)
plt.show()
```



```
In [284]: print("Decision Tree")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(tree_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(tree_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(tree_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, tree_box_forecast))))
print("MAE (test): " + str(mae(test, tree_box_forecast)))
print("sMAPE (test): " + str(smape(test, tree_box_forecast)))
```


Decision Tree

```
Observaciones ajustadas: 452
Observaciones predichas: 112
```

Test de Jarque-Bera (p-valor): 0.0010506447401931551
 Test de Ljung-Box para k = 6 (p-valor): 7.73151312209853e-06
 Test de Ljung-Box para k = 12 (p-valor): 1.3985418477241946e-06

RMSE (test): 36.49420211710823
 MAE (test): 22.4493403519271
 sMAPE (test): 52.483487500997285

```
In [285]: forecast_plot(test, tree_box_forecast)
plt.show()
```



Random Forest

```
In [286]: best_score, best_tree = float("inf"), None

n_estimators = [50, 100, 200, 500]
depths = np.linspace(1, 61, 20)
max_features = [0.8, 1]

for n in n_estimators:
    for depth in depths:
        for features in max_features:
            rf = RandomForestRegressor(n_estimators=n,
                                      max_depth=depth,
                                      max_features=features,
                                      random_state=seed)

            rf_model1 = rf.fit(X_train1, y_train1)
```

```

rf_forecast = iter_forecast(rf_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:,-1], rf_forecast))

rf_model2 = rf.fit(X_val1, y_val1)
rf_forecast = iter_forecast(rf_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], rf_forecast))

rf_model3 = rf.fit(X_val2, y_val2)
rf_forecast = iter_forecast(rf_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], rf_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_rf = cvmse, rf.get_params()

print('Best RF: n_estimators: %s max_depth: %s max_features :%s | CV RMSE = %.3f' %
      (best_rf['n_estimators'], best_rf['max_depth'], best_rf['max_features'], best_score))

Best RF: n_estimators: 200 max_depth: 10.473684210526315 max_features :1 | CV RMSE = 0.085

In [287]: rf_hyp = {"max_depth": 10.47,
                  "n_estimators": 200,
                  "max_features": 1}

In [288]: rf = RandomForestRegressor(max_depth=rf_hyp["max_depth"],
                                    n_estimators=rf_hyp["n_estimators"],
                                    max_features=rf_hyp["max_features"],
                                    random_state=seed)

rf_model = rf.fit(X_train, y_train)
rf_fitted = pd.Series(rf_model.predict(X_train))
rf_fitted
rf_fitted.index = X_train.index
rf_resid = y_train - rf_fitted
rf_resid = (rf_resid - rf_resid.mean()) / rf_resid.std()

rf_jb_test = jb(rf_resid)
rf_lb_test = lb(rf_resid)

rf_forecast = iter_forecast(rf, y_train, 12, test)

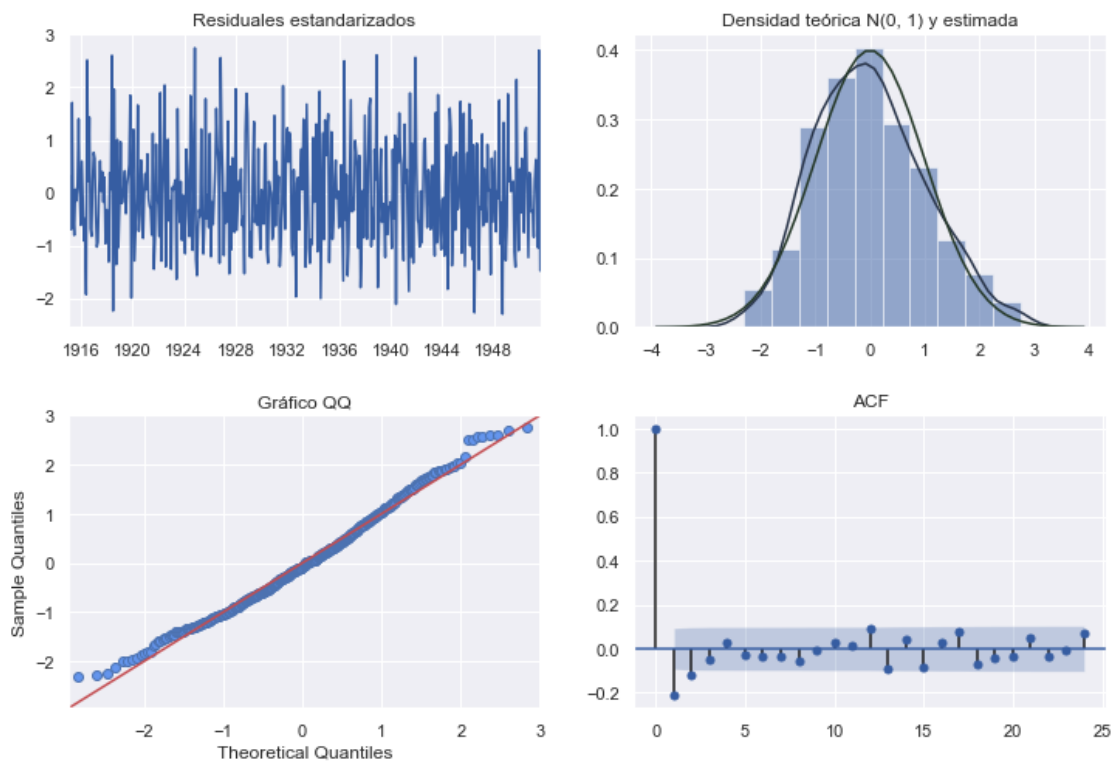
In [289]: rf_undiff_forecast = []
last_obser = bc_deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + rf_forecast[i]
    rf_undiff_forecast.append(yhat)
    last_obser = yhat

```

```
rf_undiff_forecast = pd.Series(rf_undiff_forecast)
rf_undiff_forecast.index = test.index
```

```
In [290]: rf_seas_forecast = rf_undiff_forecast + bc_seasonality_forecast
rf_box_forecast = (bc_param * rf_seas_forecast + 1) ** (1 / bc_param)
```

```
In [291]: resid_diag(rf_resid)
plt.show()
```



```
In [292]: print("Random Forest")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(rf_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(rf_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(rf_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, rf_box_forecast))))
print("MAE (test): " + str(mae(test, rf_box_forecast)))
print("sMAPE (test): " + str(smape(test, rf_box_forecast)))
```

Random Forest

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 0.011941226814951477

Test de Ljung-Box para k = 6 (p-valor): 0.00015047040565678517


Test de Ljung-Box para k = 12 (p-valor): 0.0002746332719588352

RMSE (test): 33.32998445529877

MAE (test): 20.927560058562502

sMAPE (test): 43.07766738416892

```
In [293]: forecast_plot(test, rf_box_forecast)
plt.show()
```



ElasticNet

```
In [294]: best_score, best_tree = float("inf"), None

ratio = [0.1, 0.5, 0.7, 0.9, 0.95, 0.99, 1]
alpha = [0, 0.01, 0.1, 0.5, 1, 5, 10, 50, 100]

for rat in ratio:
    for alp in alpha:
        en = ElasticNet(l1_ratio=rat, alpha=alp, random_state=seed)

        en_model1 = en.fit(X_train1, y_train1)
```



```

en_forecast = iter_forecast(en_model1, y_train1, 12, validation1)
rmse1 = sqrt(mse(validation1.iloc[:,-1], en_forecast))

en_model2 = en.fit(X_val1, y_val1)
en_forecast = iter_forecast(en_model2, y_val1, 12, validation2)
rmse2 = sqrt(mse(validation2.iloc[:,-1], en_forecast))

en_model3 = en.fit(X_val2, y_val2)
en_forecast = iter_forecast(en_model3, y_val2, 12, validation3)
rmse3 = sqrt(mse(validation3.iloc[:,-1], en_forecast))

cvmse = mean([rmse1, rmse2, rmse3])
if cvmse < best_score:
    best_score, best_en = cvmse, en.get_params()

print('Best EN: ratio: %s alpha: %s | CV RMSE = %.3f' %
      (best_en['l1_ratio'], best_en['alpha'], best_score))

```

Best EN: ratio: 0.1 alpha: 0.1 | CV RMSE = 0.085

```

In [295]: en_hyp = {"l1_ratio": 0.1,
                  "alpha": 0.1}

```

```

In [296]: en = ElasticNet(l1_ratio=en_hyp["l1_ratio"], alpha=en_hyp["alpha"],
                        random_state=seed)
en_model = en.fit(X_train, y_train)
en_fitted = pd.Series(en_model.predict(X_train))
en_fitted
en_fitted.index = X_train.index
en_resid = y_train - en_fitted
en_resid = (en_resid - en_resid.mean()) / en_resid.std()

en_jb_test = jb(en_resid)
en_lb_test = lb(en_resid)

en_forecast = iter_forecast(en, y_train, 12, test)

```

```

In [297]: en_undiff_forecast = []
last_obser = bc_deseas.iloc[-1].value
for i in range(len(test)):
    yhat = last_obser + en_forecast[i]
    en_undiff_forecast.append(yhat)
    last_obser = yhat

```

```

en_undiff_forecast = pd.Series(en_undiff_forecast)
en_undiff_forecast.index = test.index

```

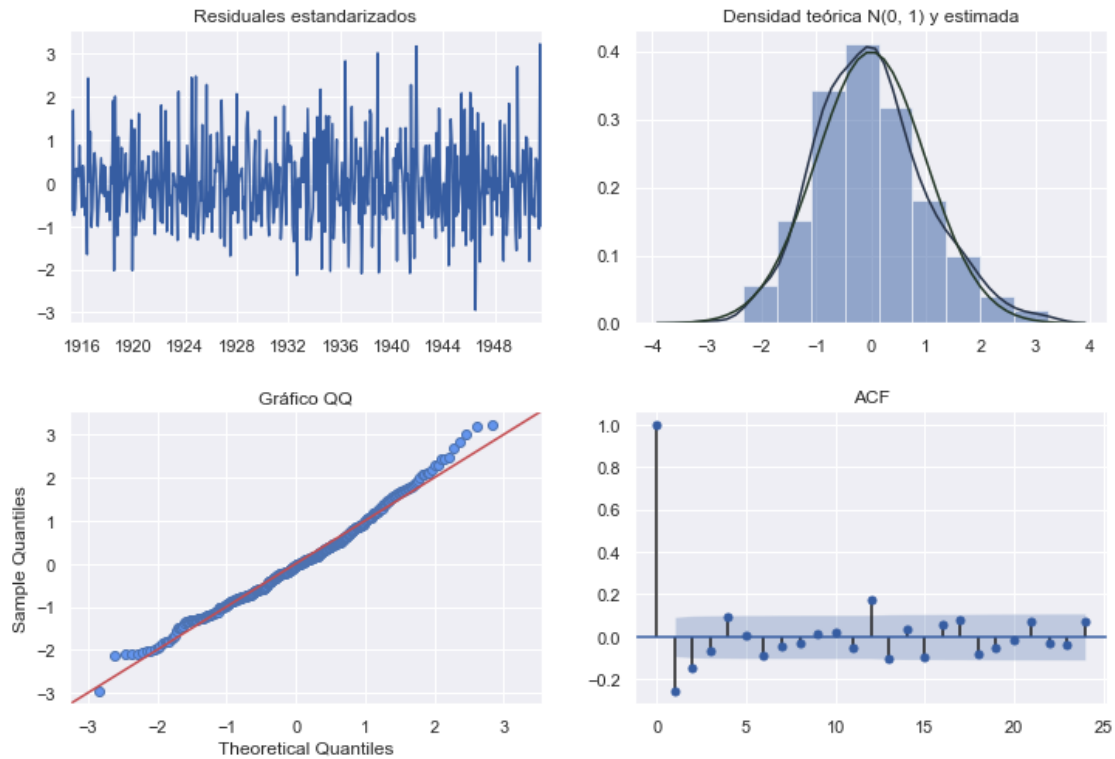
```

In [298]: en_seas_forecast = en_undiff_forecast + bc_seasonality_forecast
en_box_forecast = (bc_param * en_seas_forecast + 1) ** (1 / bc_param)

```

```
In [299]: en_box_forecast = (bc_param * en_seas_forecast + 1) ** (1 / bc_param)
```

```
In [300]: resid_diag(en_resid)
plt.show()
```



```
In [301]: print("ElasticNet")
print("")
print("Observaciones ajustadas: " + str(len(train)))
print("Observaciones predichas: " + str(len(test)))
print("")
print("Test de Jarque-Bera (p-valor): " + str(en_jb_test[1]))
print("Test de Ljung-Box para k = 6 (p-valor): " +
      str(en_lb_test[1][6]))
print("Test de Ljung-Box para k = 12 (p-valor): " +
      str(en_lb_test[1][12]))
print("")
print("RMSE (test): " + str(sqrt(mse(test, en_box_forecast))))
print("MAE (test): " + str(mae(test, en_box_forecast)))
print("sMAPE (test): " + str(smape(test, en_box_forecast)))
```

ElasticNet

Observaciones ajustadas: 452

Observaciones predichas: 112

Test de Jarque-Bera (p-valor): 0.001788964092508347

Test de Ljung-Box para $k = 6$ (p-valor): 2.3939918413298452e-08

Test de Ljung-Box para $k = 12$ (p-valor): 1.0498455598897642e-09

RMSE (test): 35.28368401389798

MAE (test): 23.133149094283

sMAPE (test): 45.15517516895144

```
In [302]: forecast_plot(test, en_box_forecast)
plt.show()
```

